


Federico Muscarà Marco Parentin Leonardo Straccali Francesco Zanasi

Modeling Infection Propagation in Populations

**High Performance Computing Project
Politecnico di Torino**



Presentation Overview



- Introduction




- Solution Description



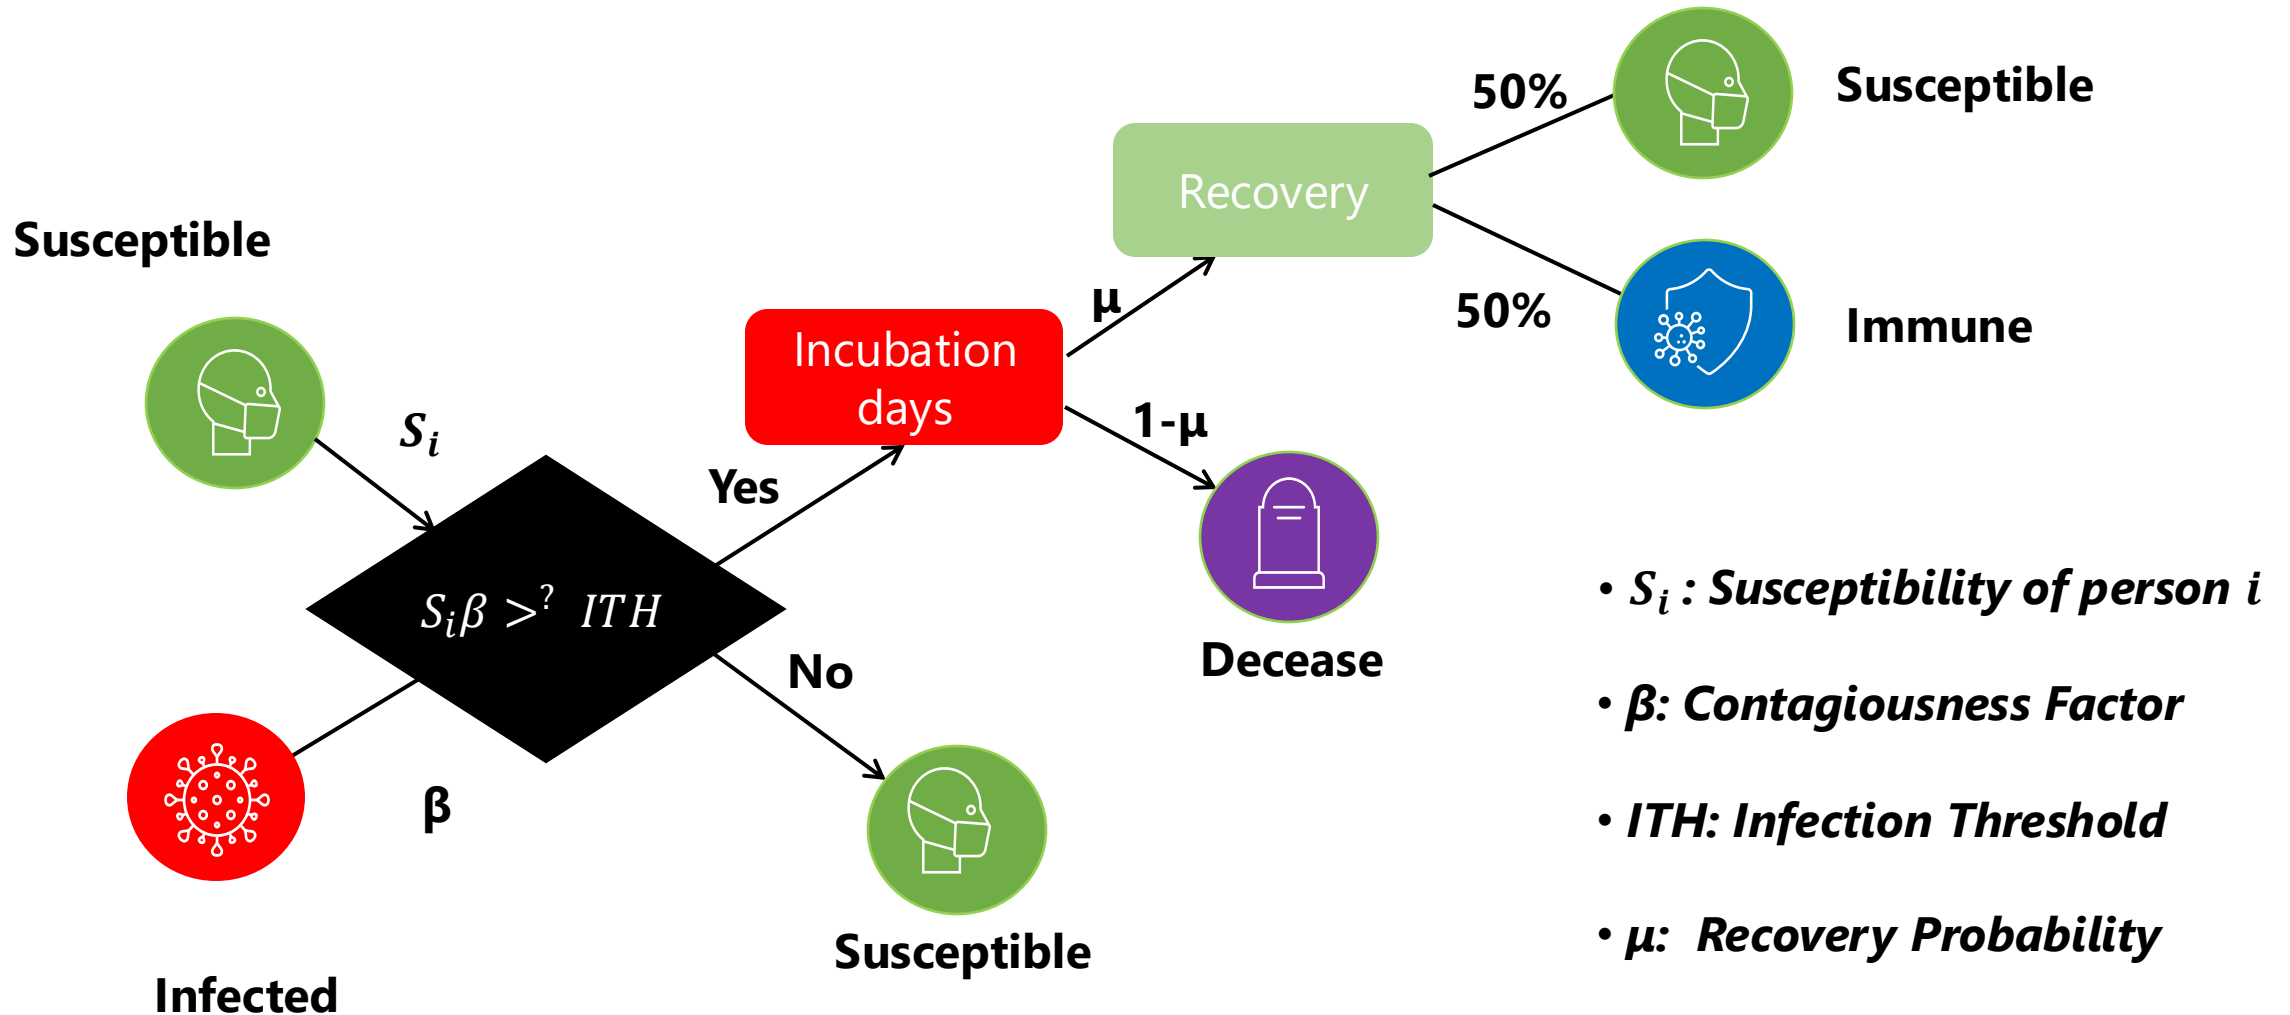
- Experimental Results



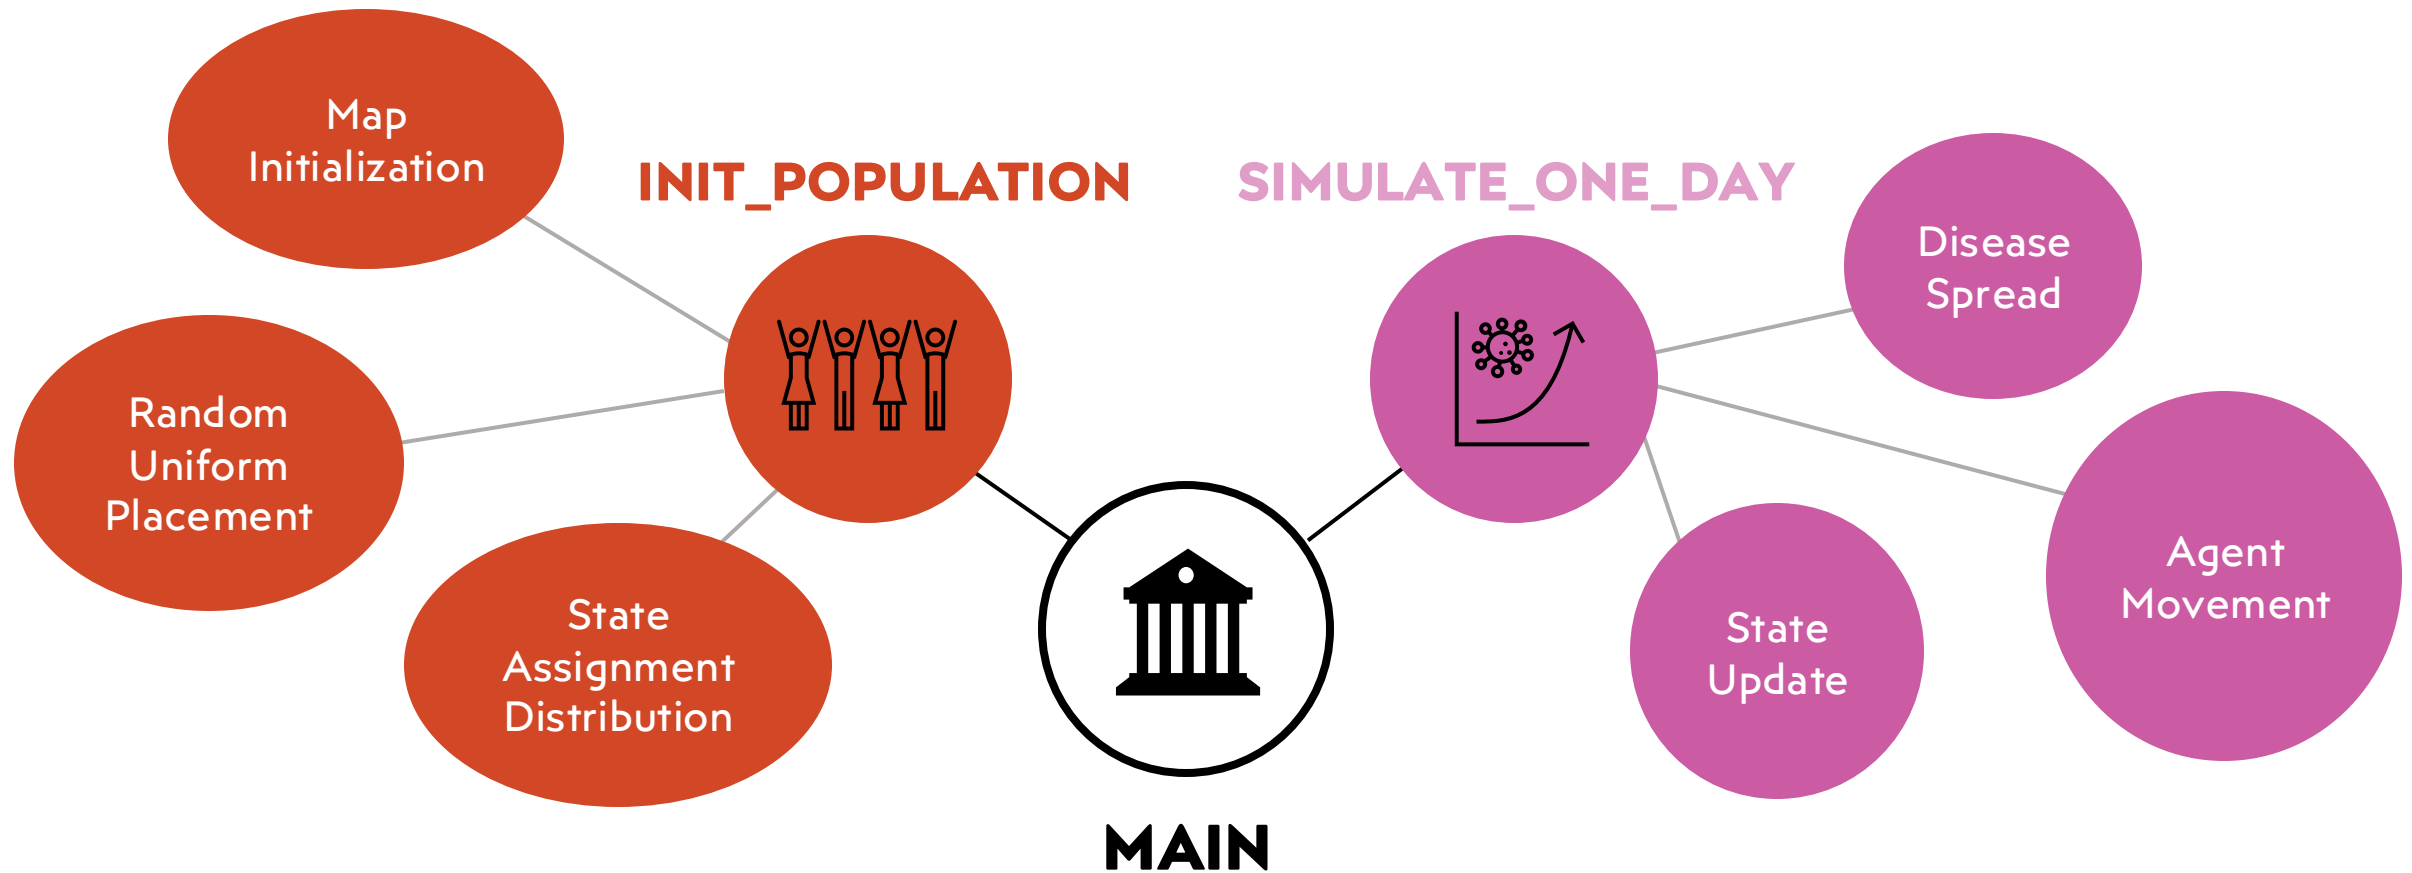
- Conclusions




Introduction: The Kermack-McKendrick (SIR) Model



Solution Description: SERIAL Version – Program Organization

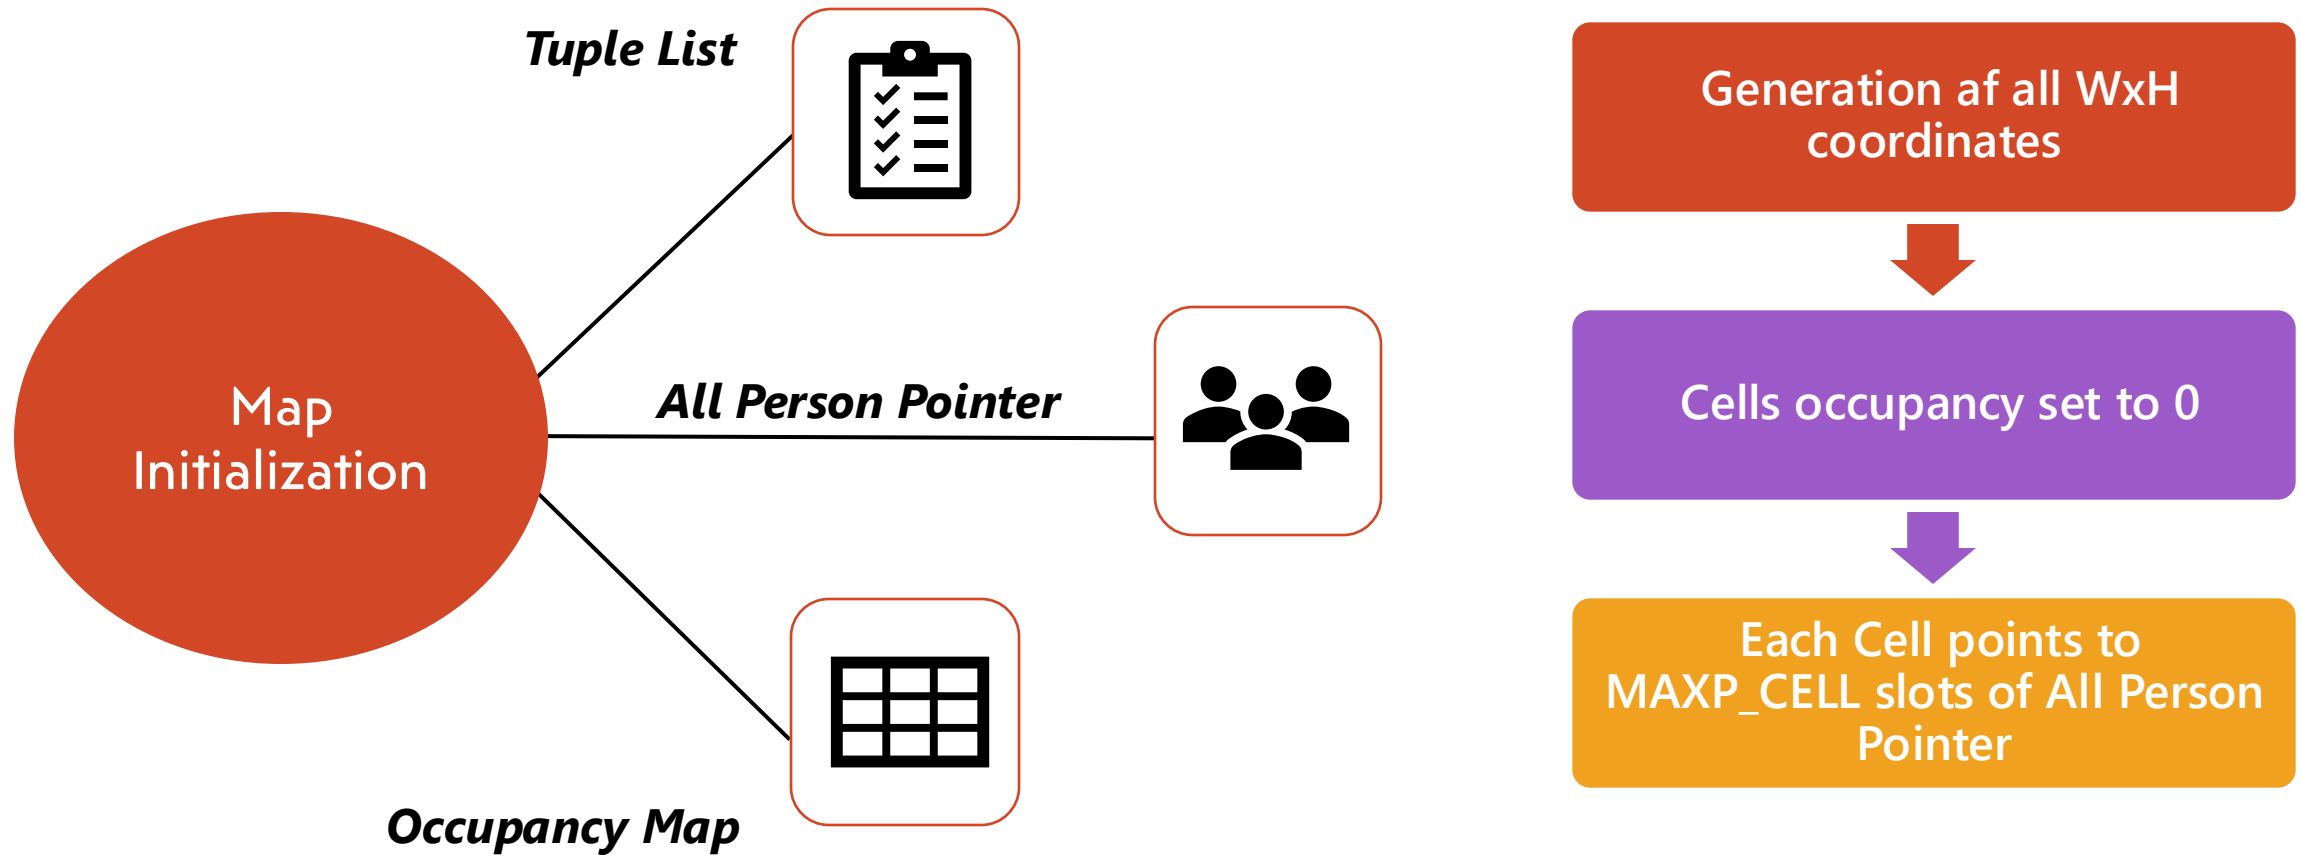


Solution Description: SERIAL Version – Population Initialization

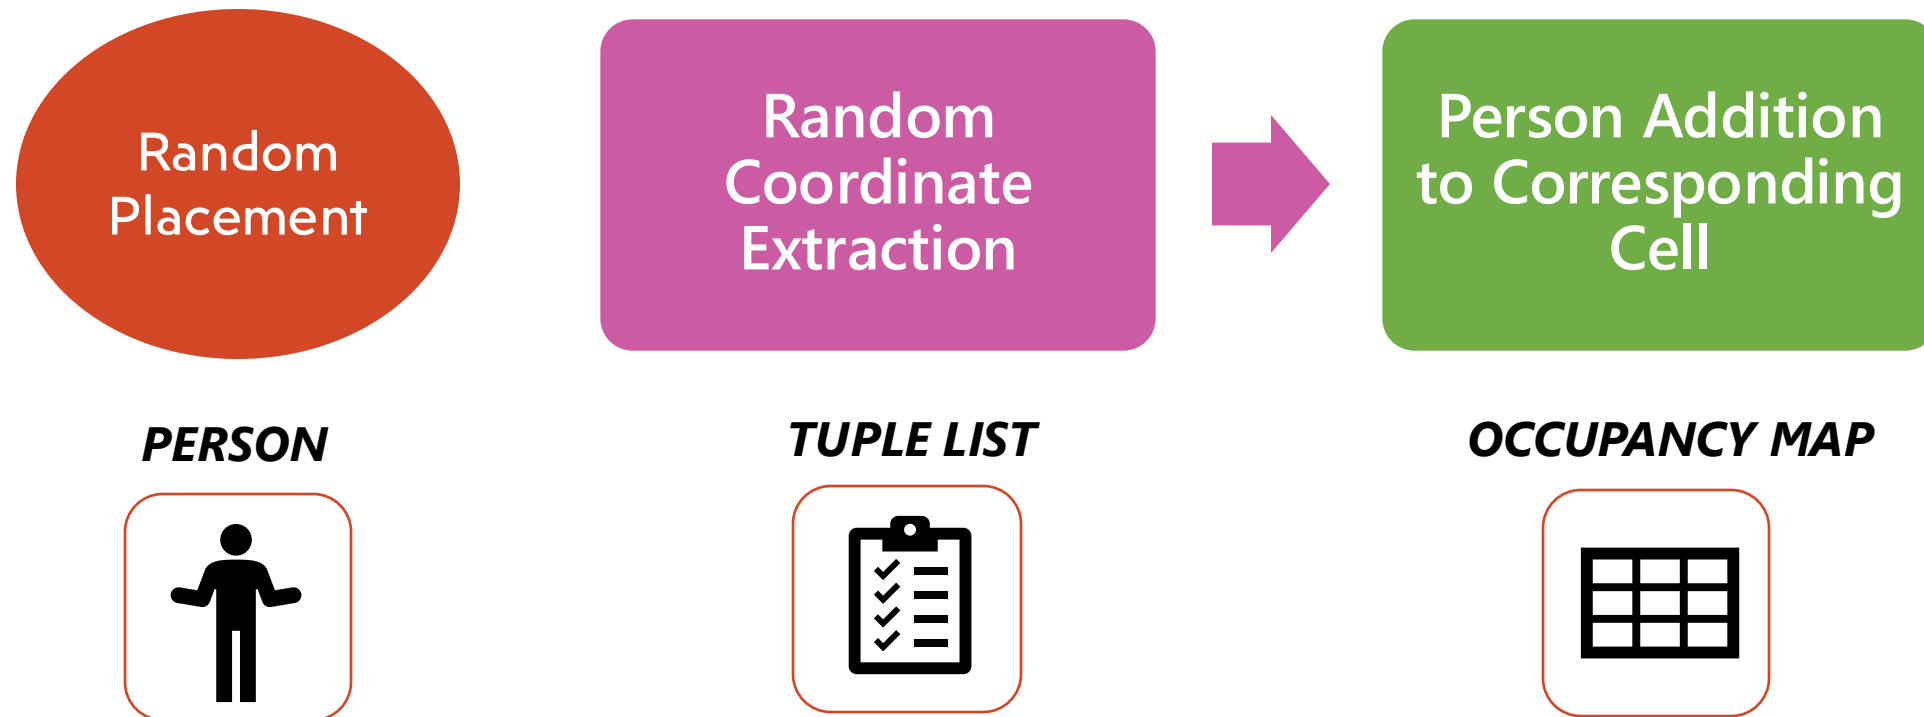


- *Each Cell contains at maximum "MAXP_CELL" individuals*

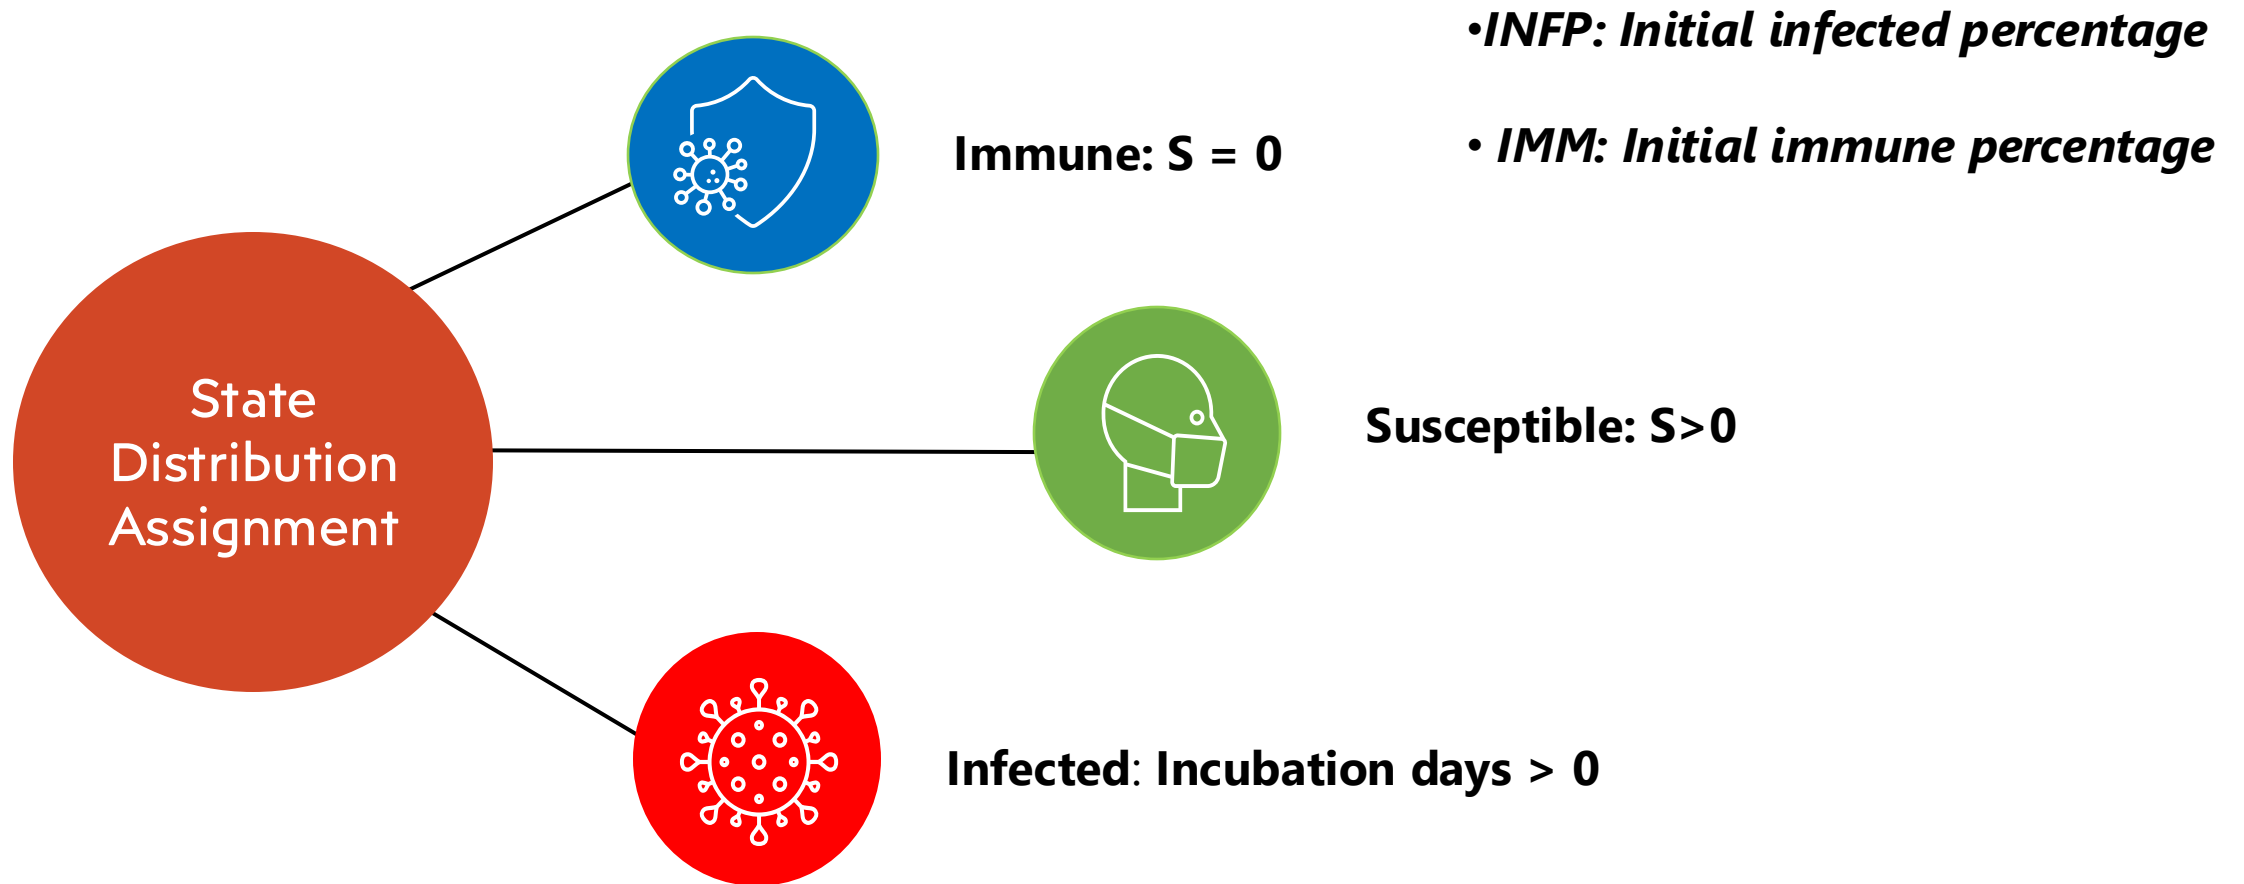
Solution Description: SERIAL Version – Population Initialization



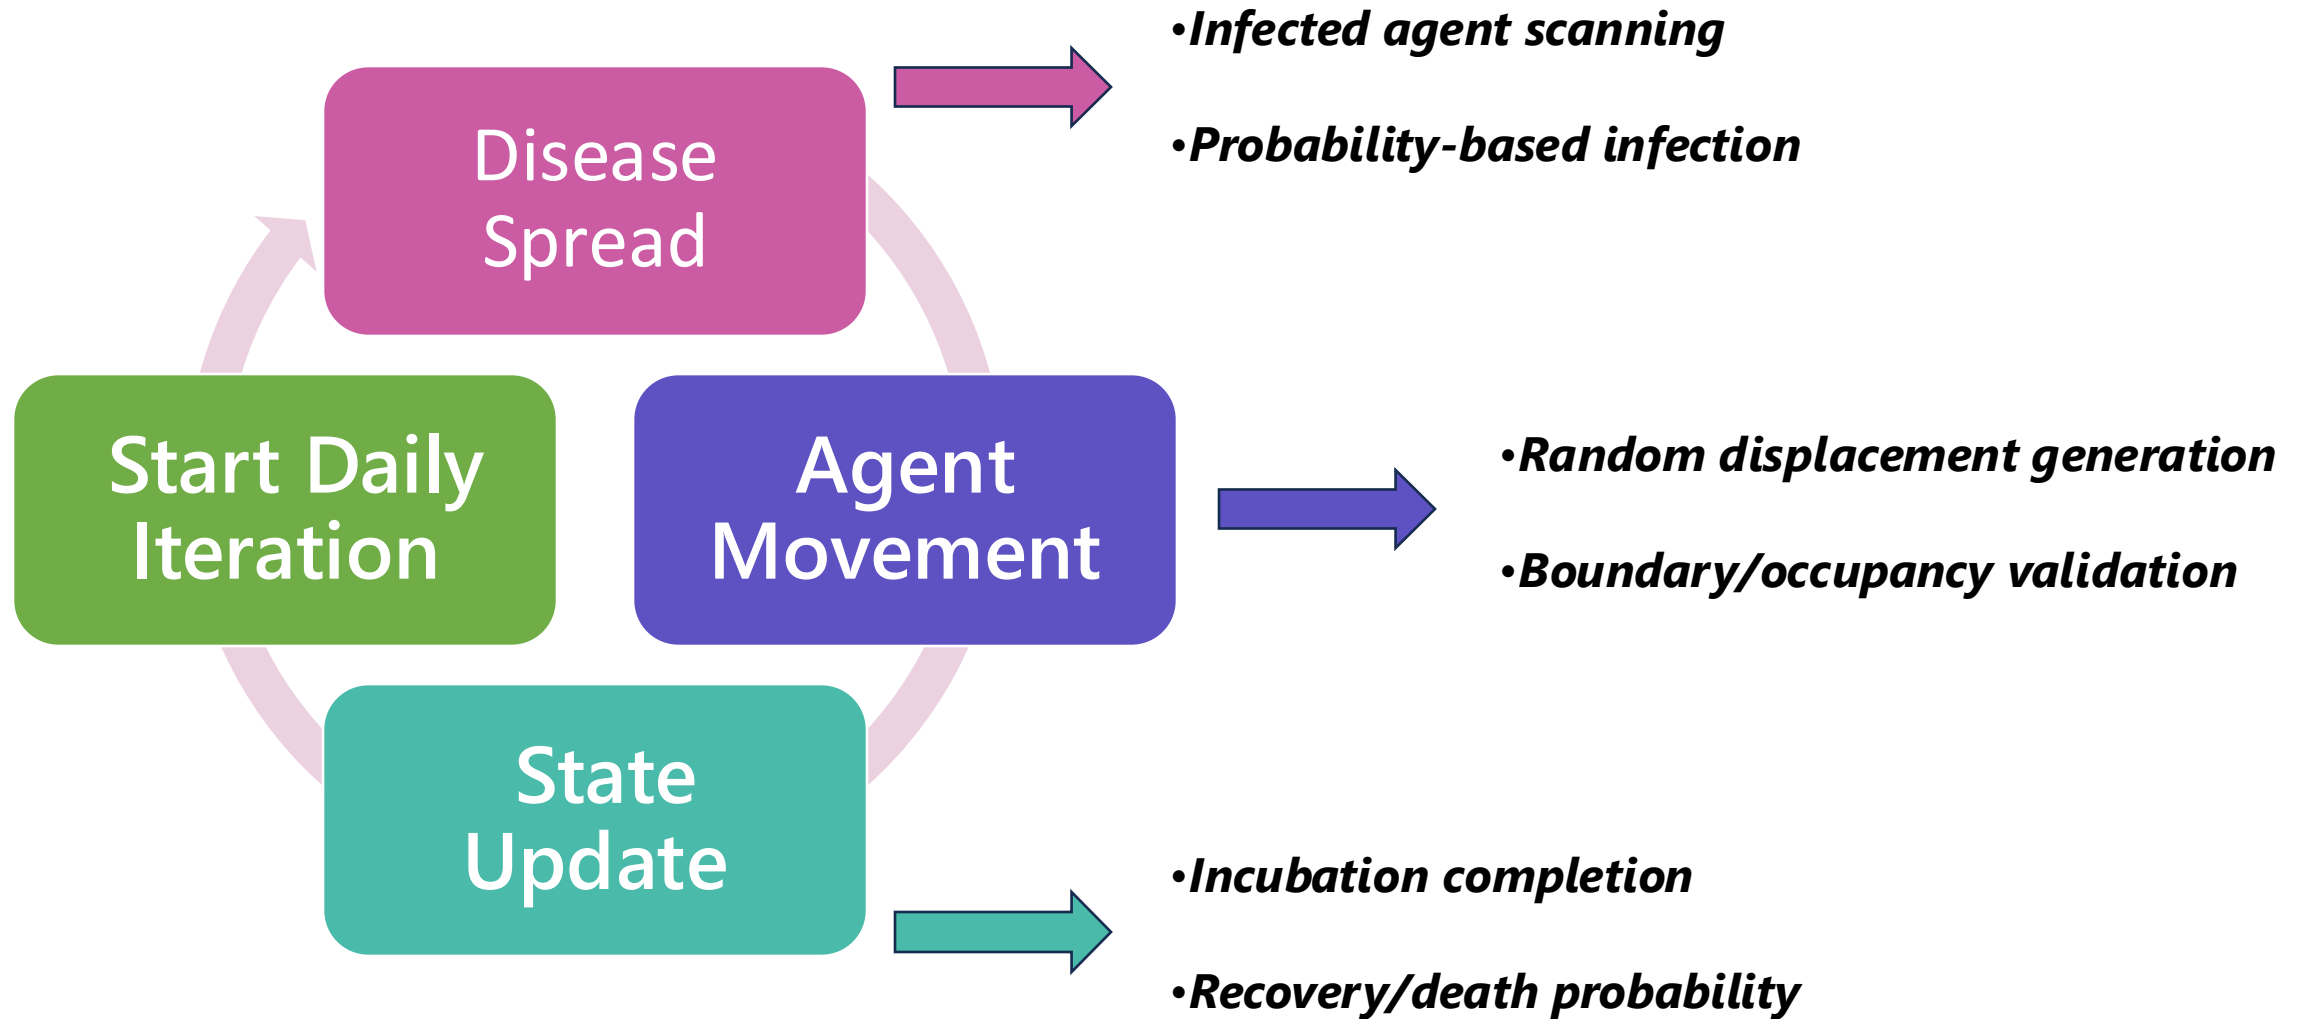
Solution Description: SERIAL Version – Population Initialization




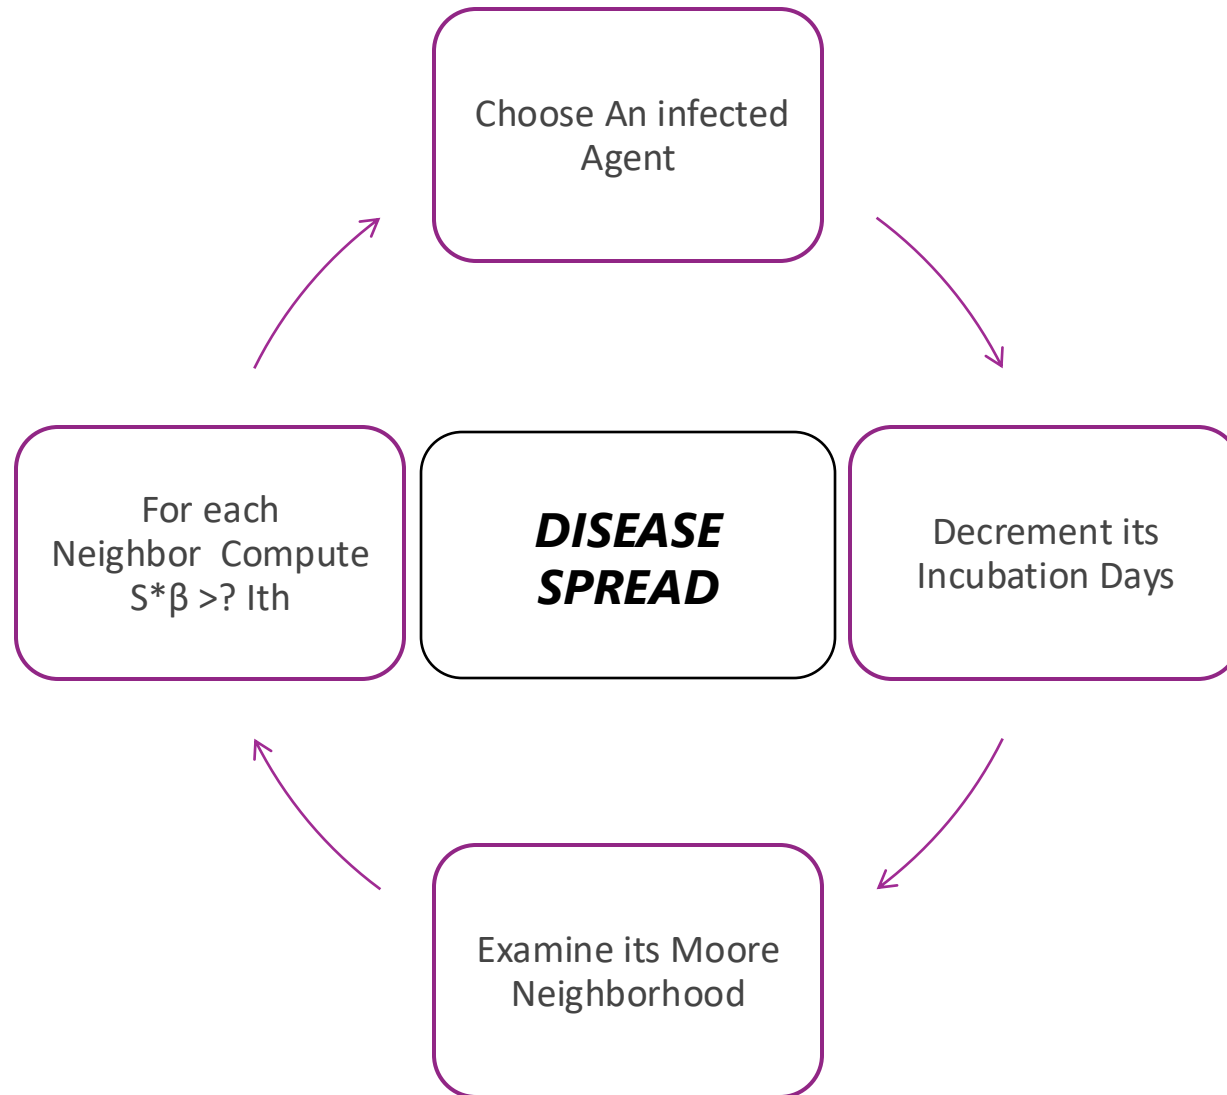
Solution Description: SERIAL Version – Population Initialization




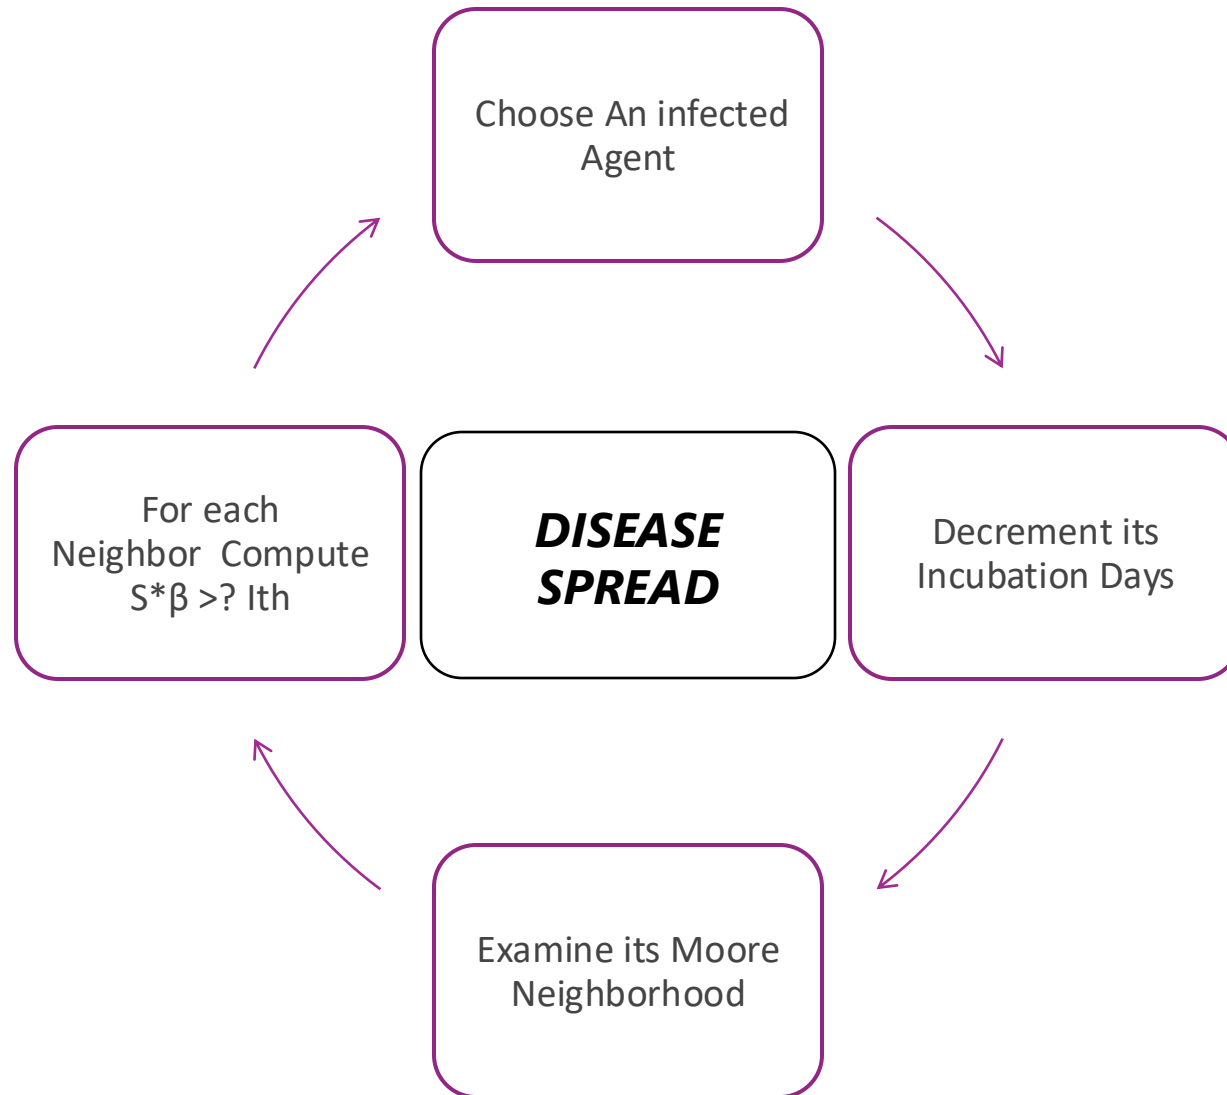
Solution Description: SERIAL Version – Daily Simulation



Solution Description: SERIAL Version – Daily Simulation





Solution Description: SERIAL Version – Daily Simulation




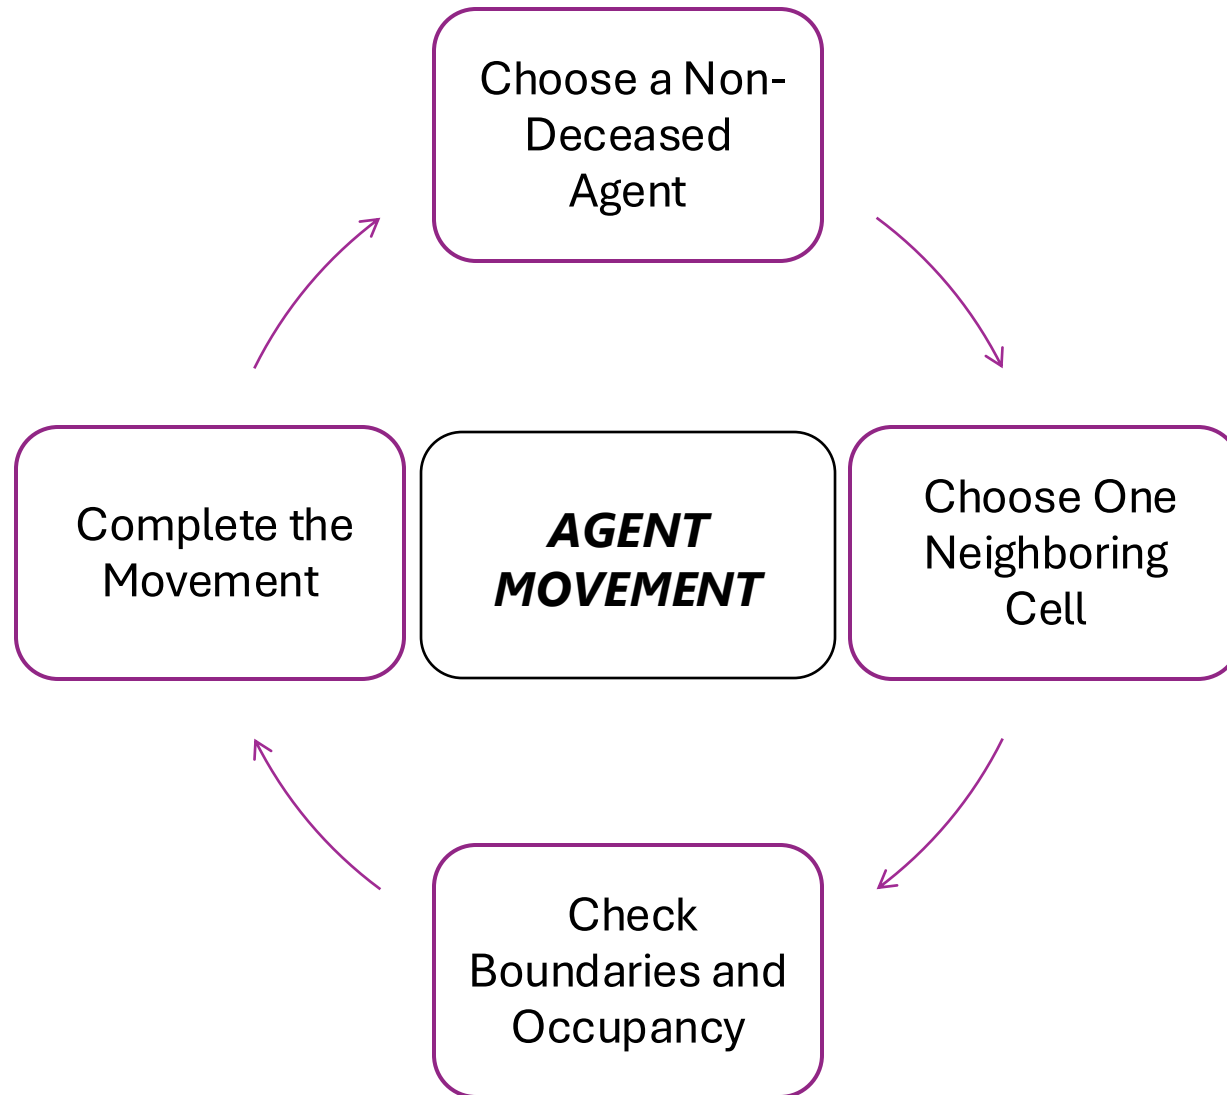
- ***Newly infected***
- ***Incubation_days = -1***




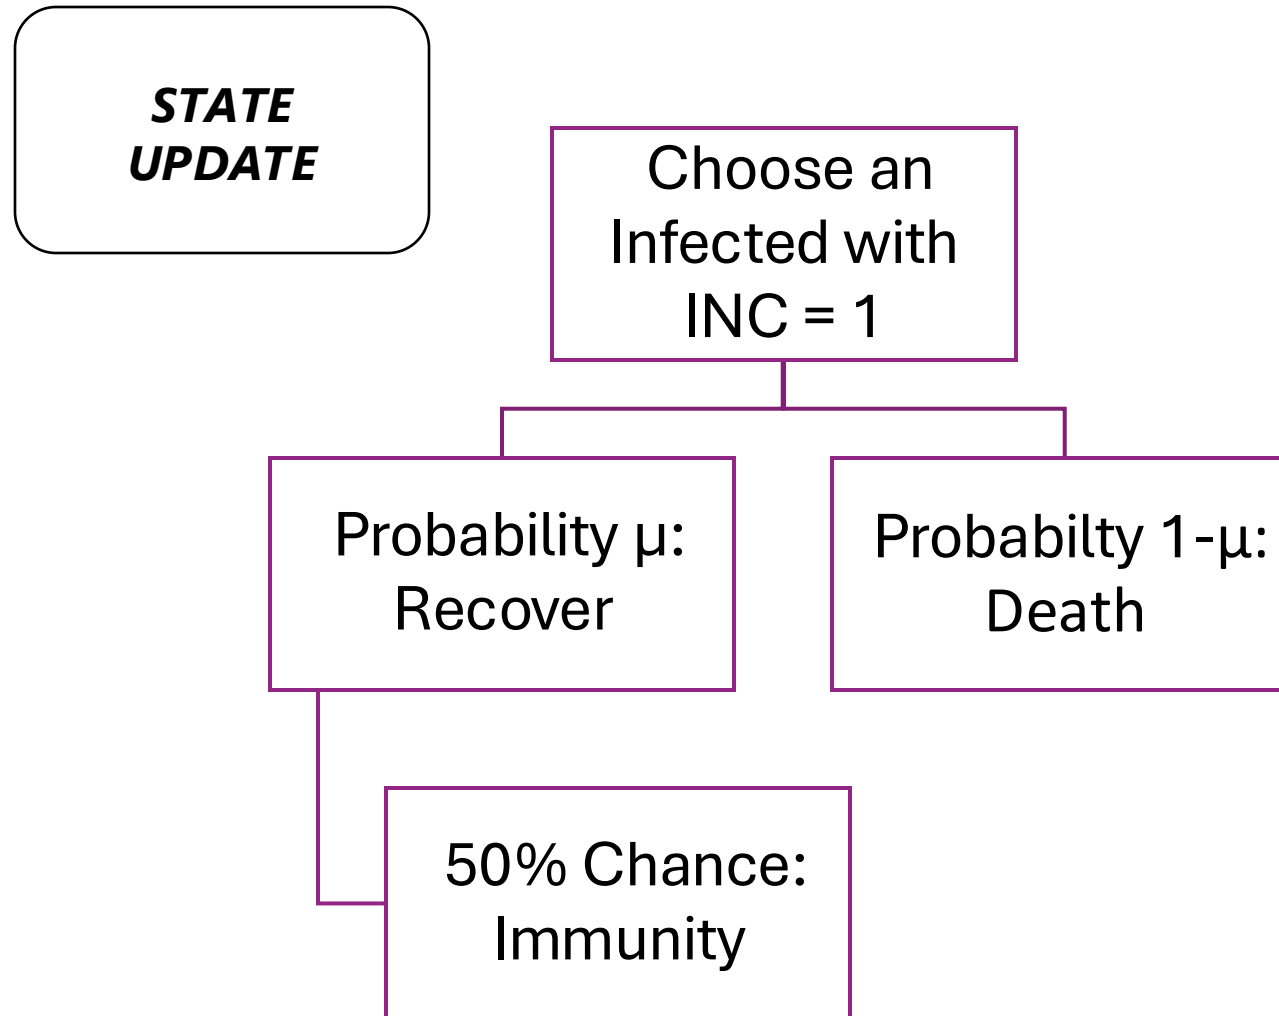
Solution Description: SERIAL Version – Daily Simulation




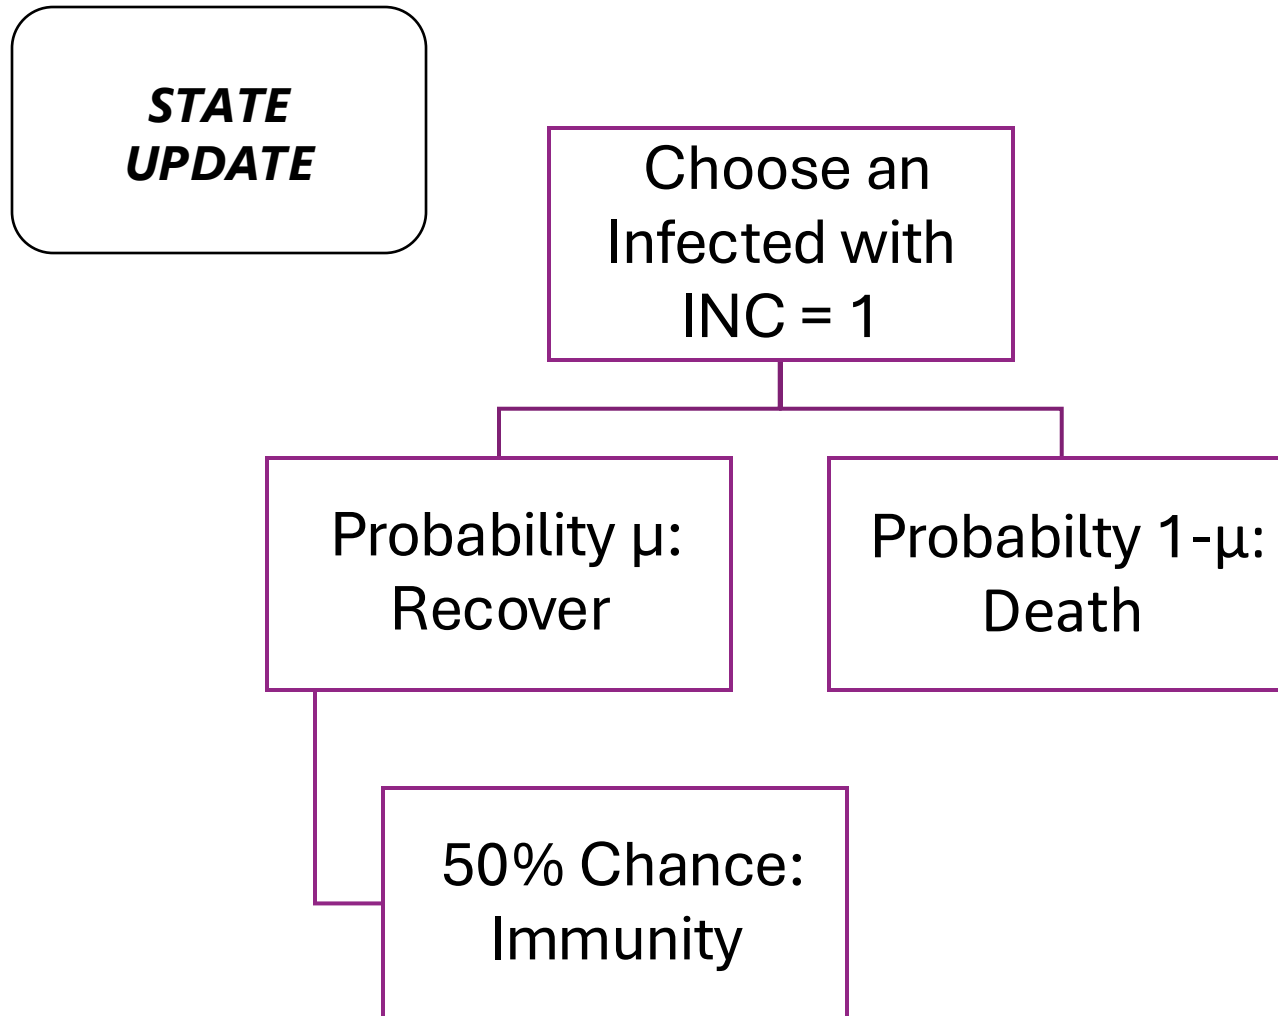
Solution Description: SERIAL Version – Daily Simulation



Solution Description: SERIAL Version – Daily Simulation



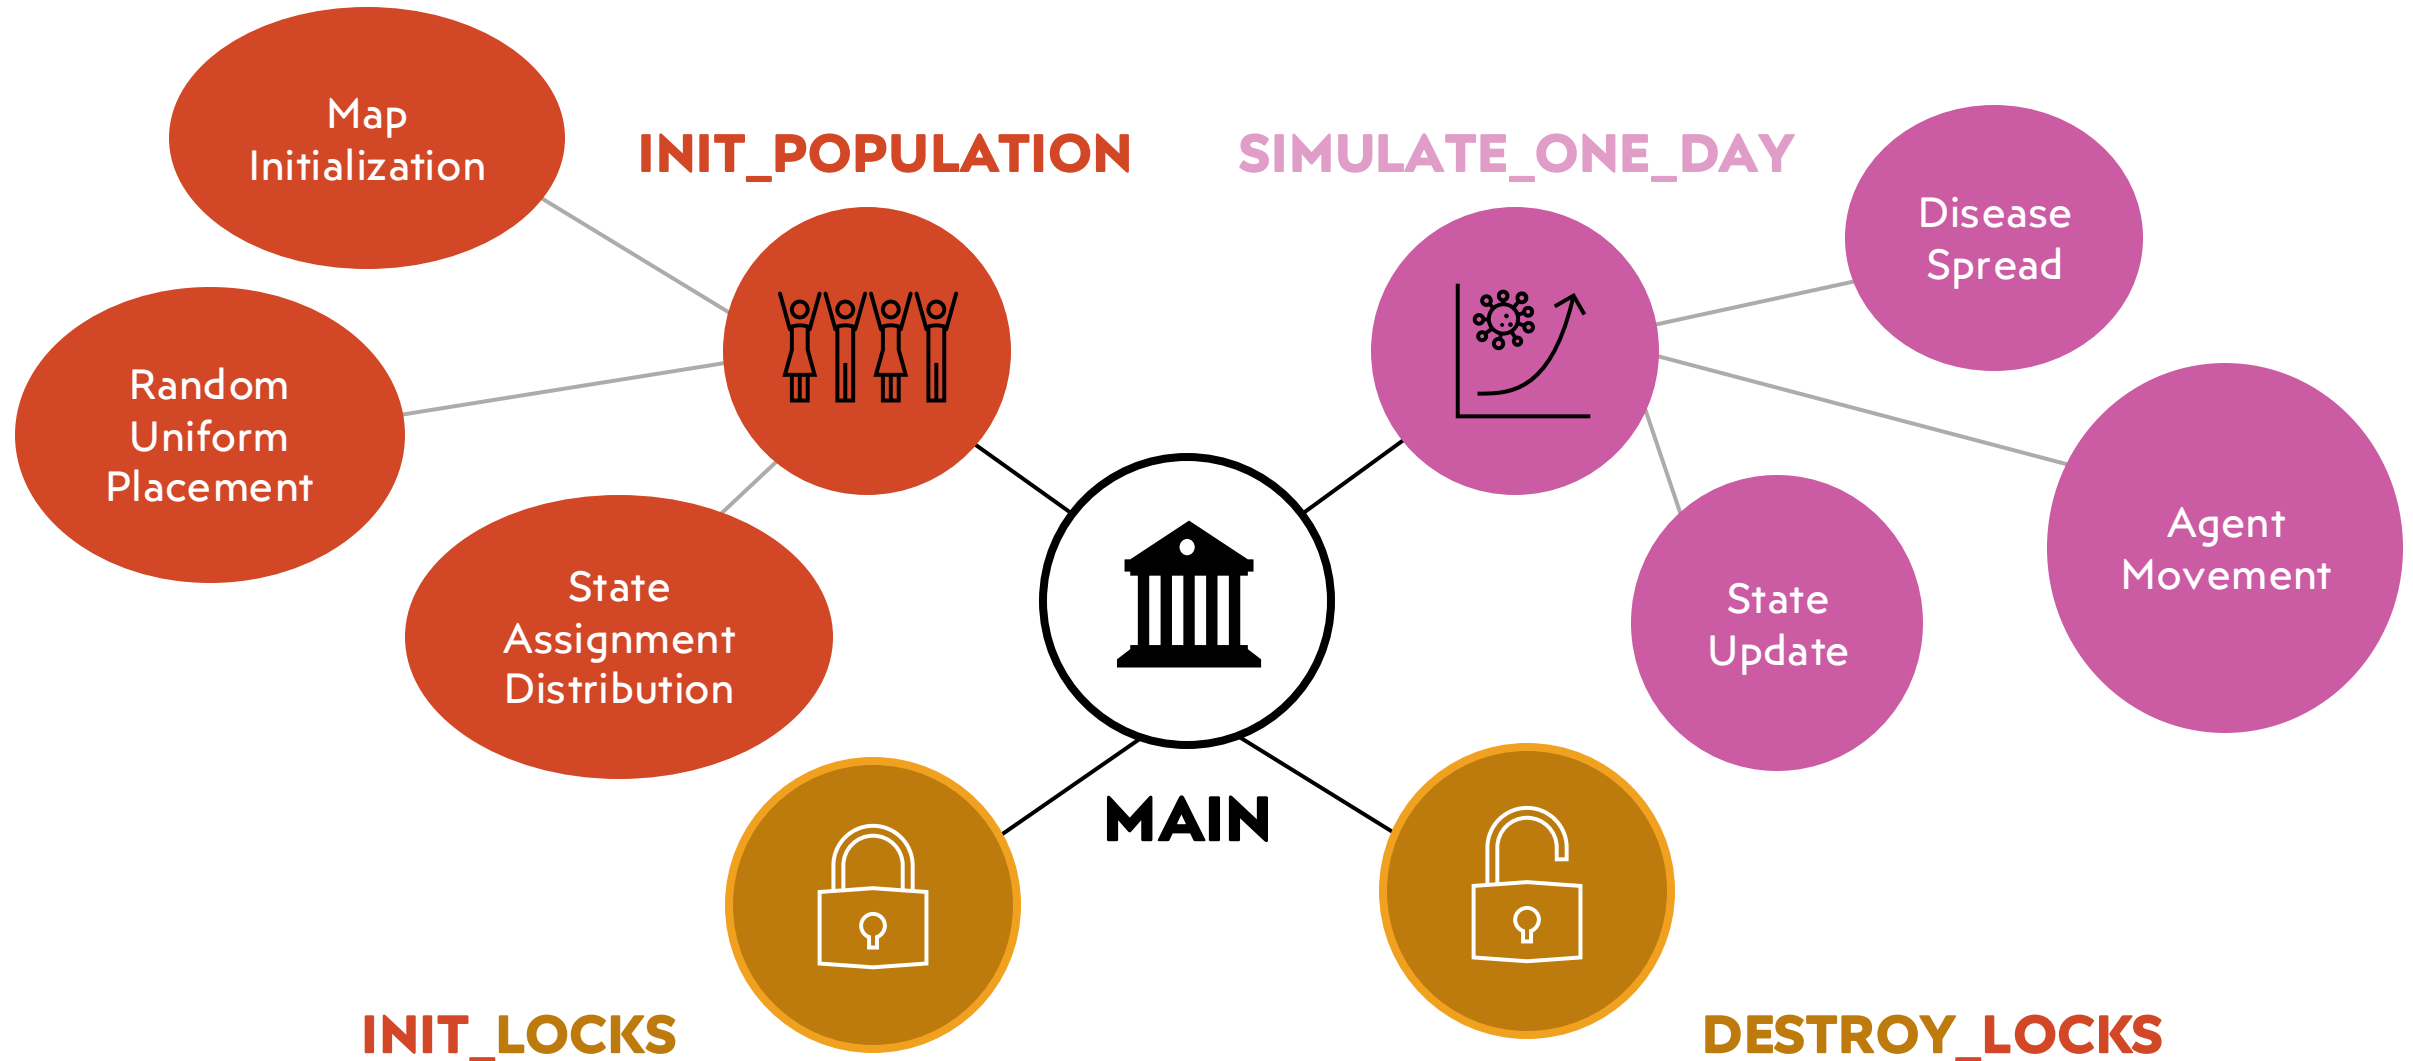
Solution Description: SERIAL Version – Daily Simulation



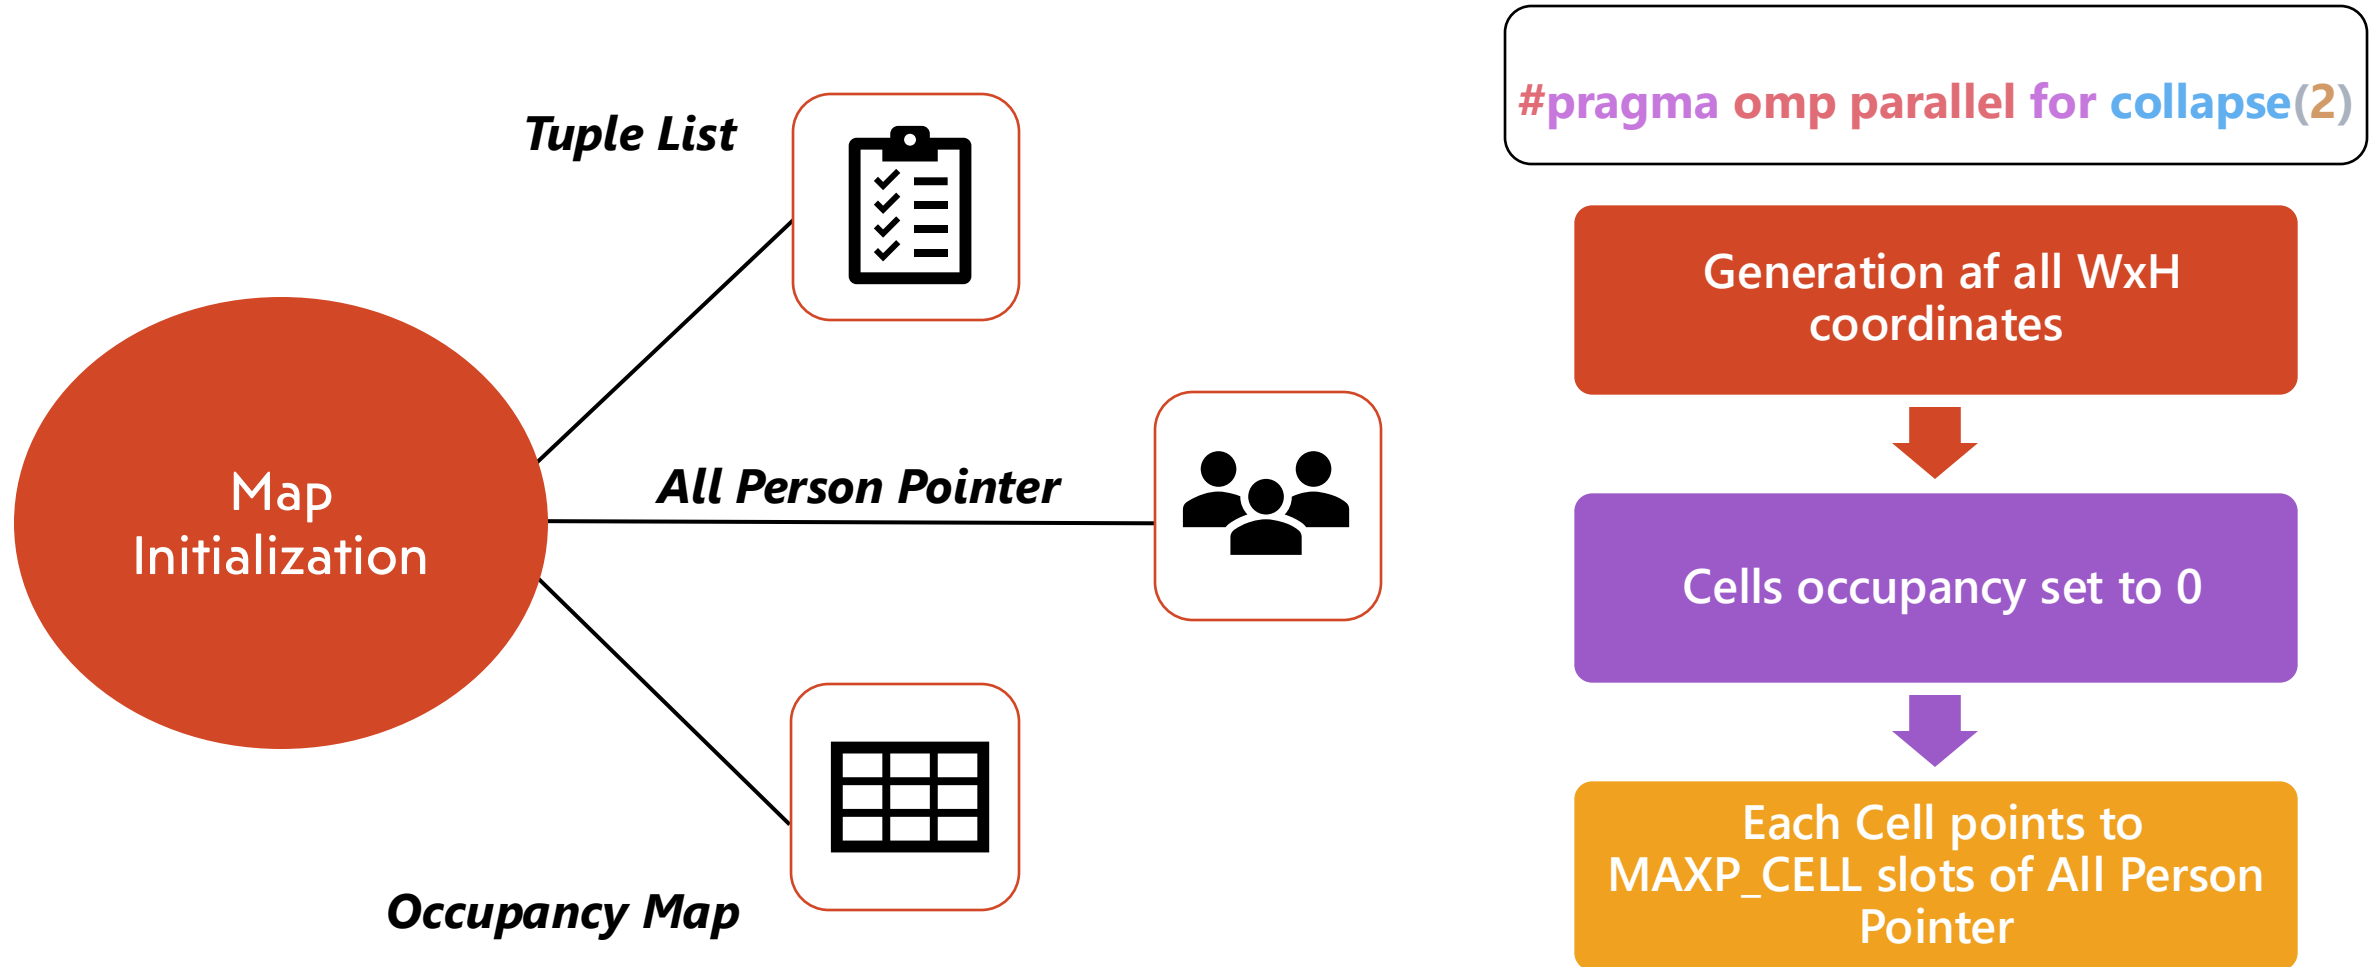
- ***Newly recovered***
- ***Incubation_days = -2***



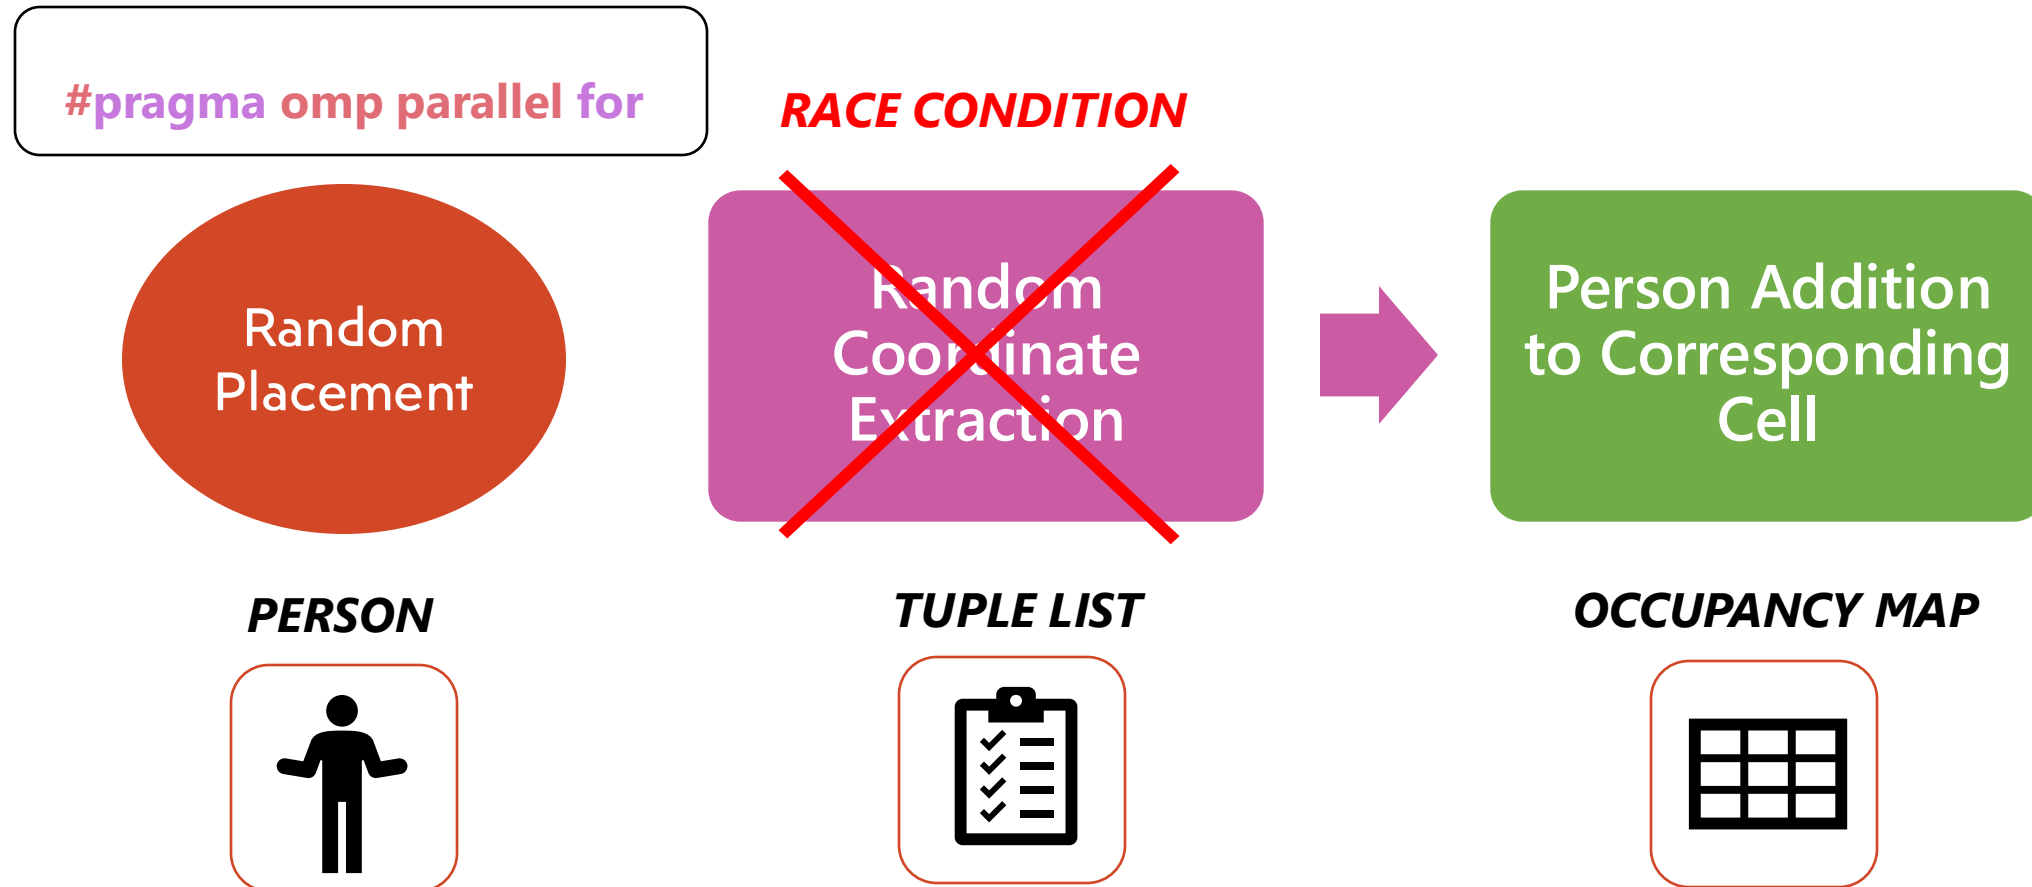
Solution Description: OpenMP Version – Program Organization



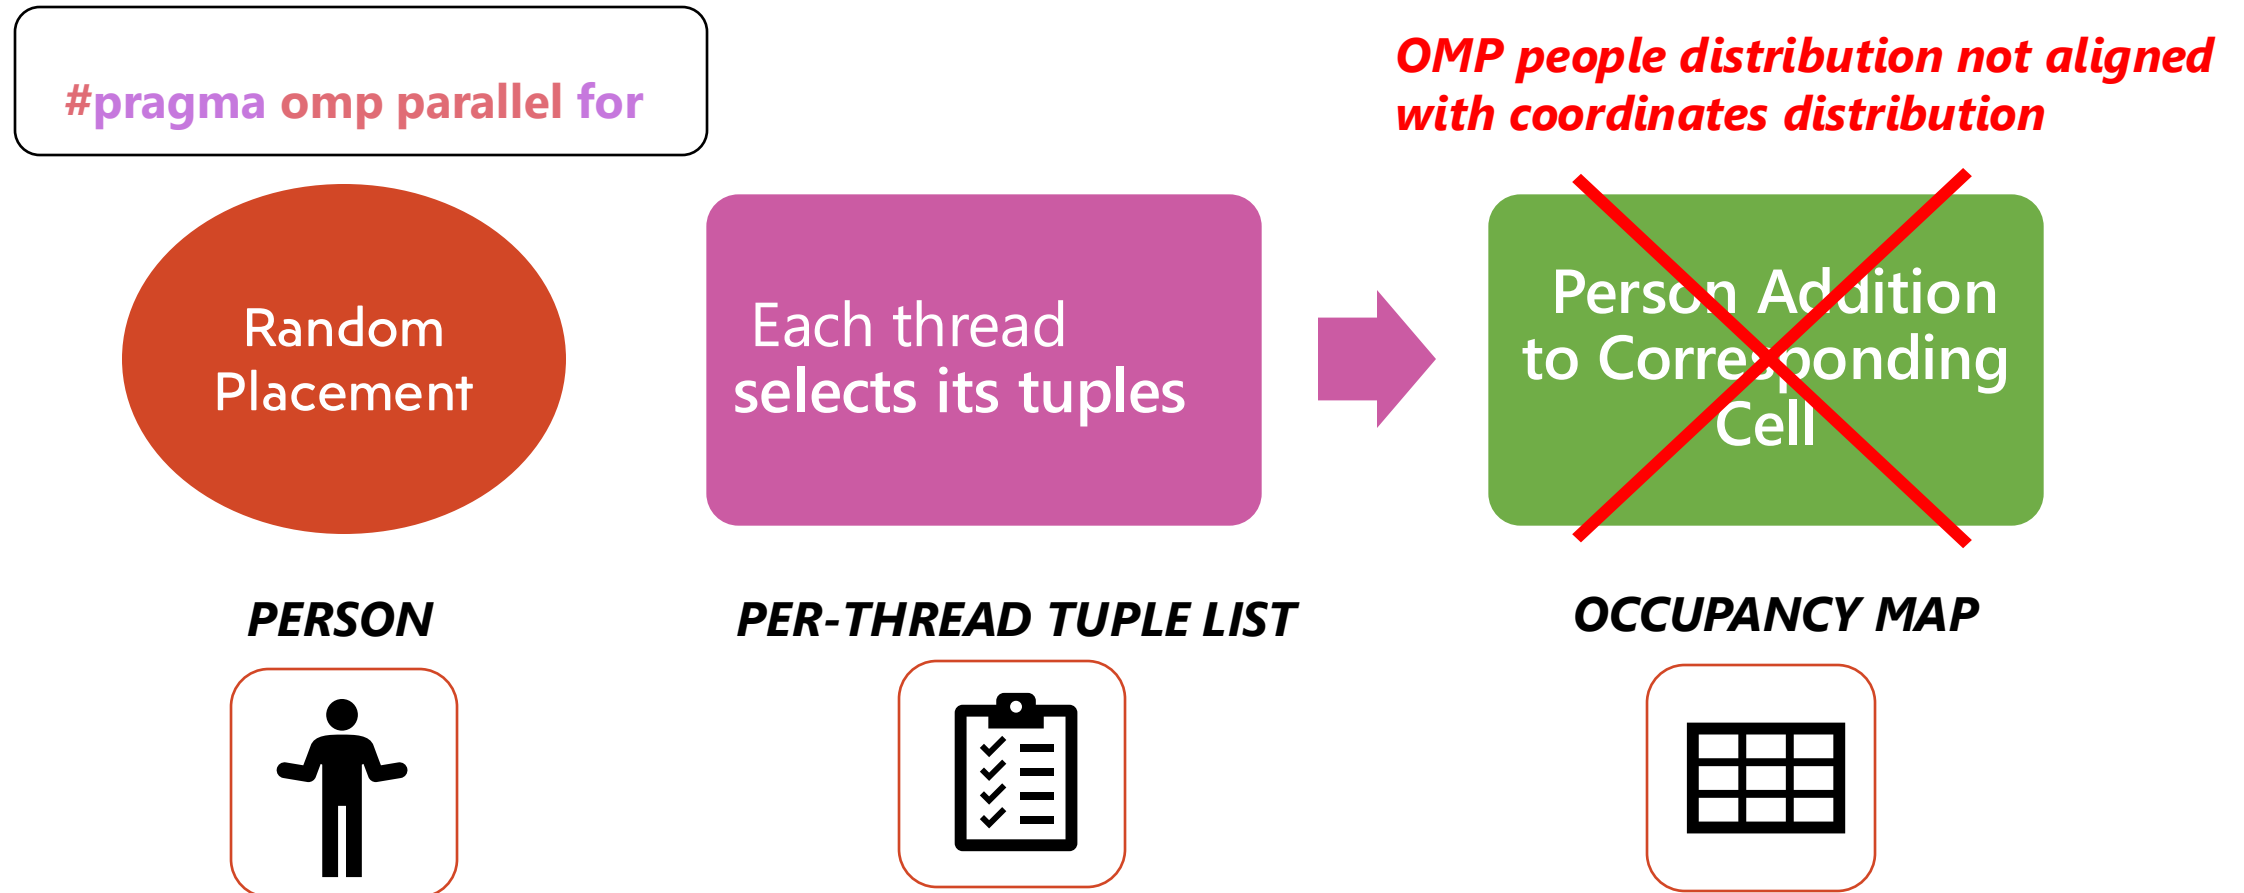
Solution Description: OpenMP Version – Population Initialization



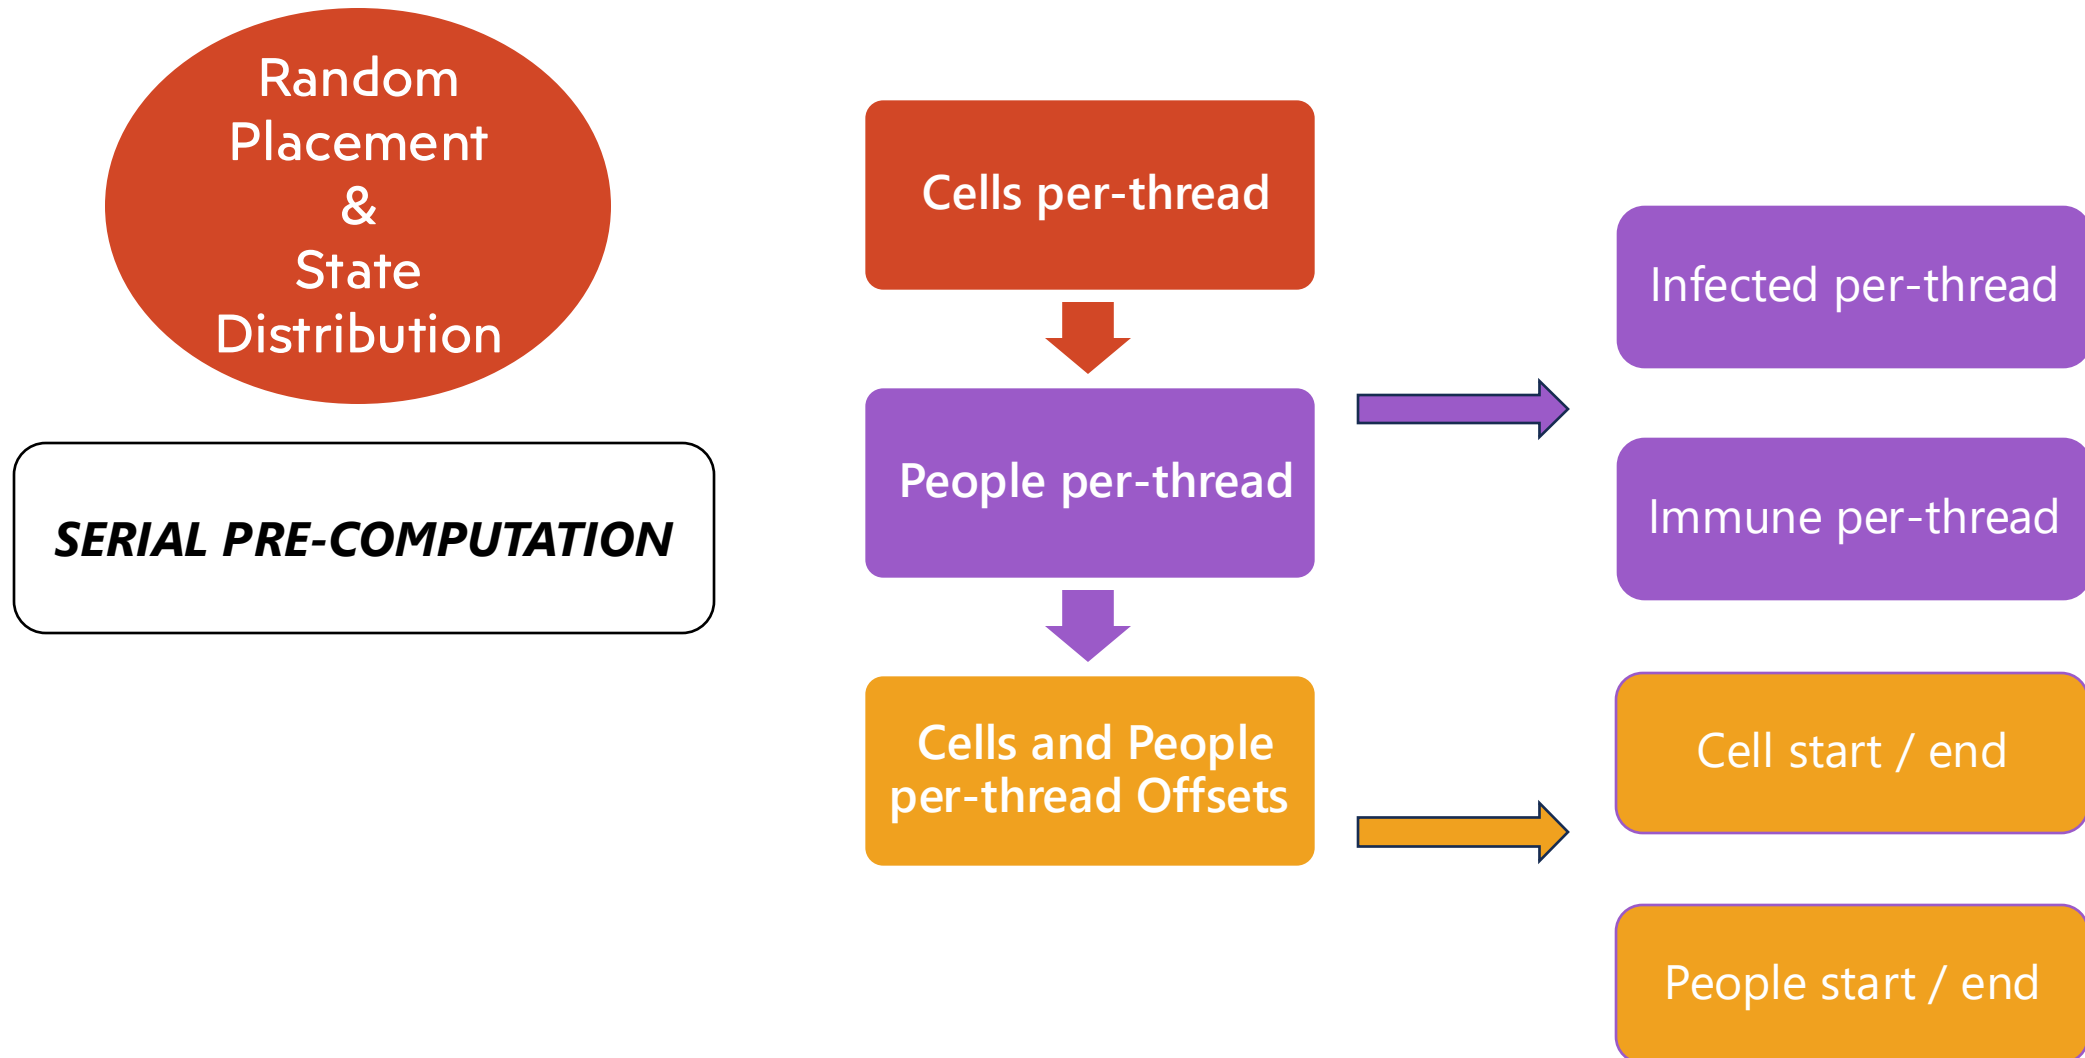
Solution Description: OpenMP Version – Population Initialization



Solution Description: OpenMP Version – Population Initialization



Solution Description: OpenMP Version – Population Initialization



Solution Description: OpenMP Version – Population Initialization

PARALLEL SECTION

Cell start / end




```
for (int i = cell_start; i < cell_end; i++)  
{  
    int x = i / H;  
    int y = i % H;  
    addTuple(local_coords, x, y);  
}
```

Random
Placement
&
State
Distribution



```
for (int i = people_start; i < people_end; i++)  
{  
    Person *p = &population[i];  
    // ...  
}
```

People start / end



TID 1



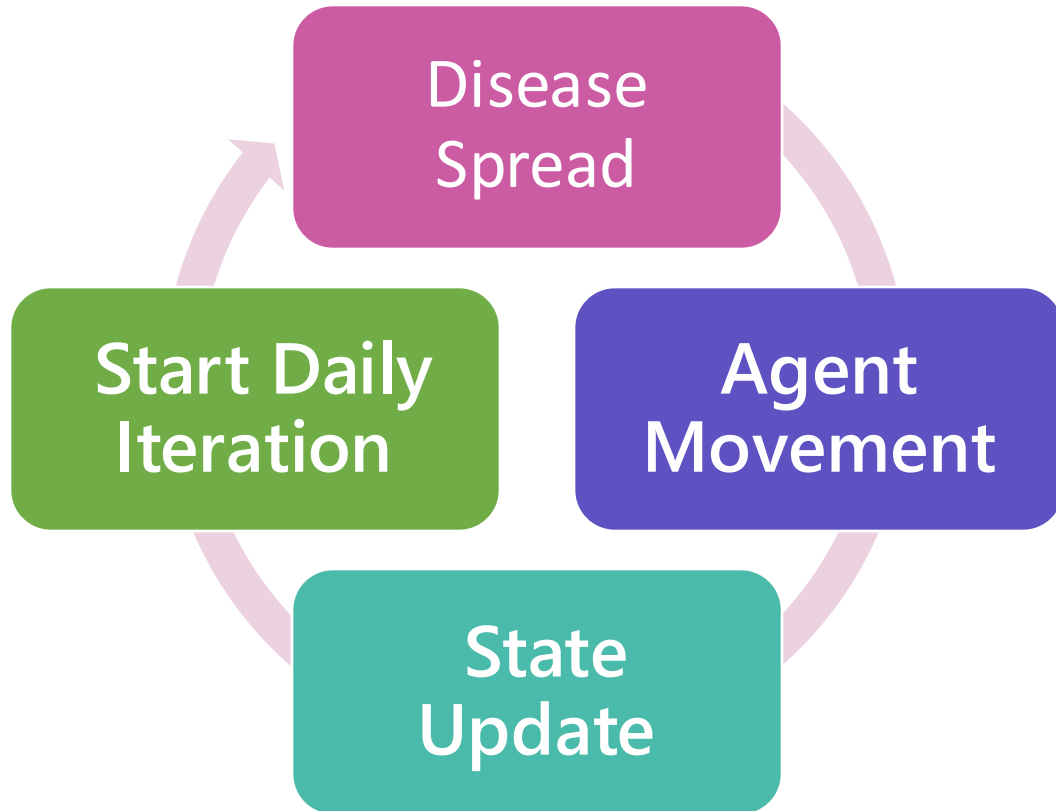
TID 2












TID 3



Solution Description: OpenMP Version – Daily Simulation





CELL LOCKS



		
		
		

- **Array of "omp_lock_t"**
- **Each lock is associated to a tuple**

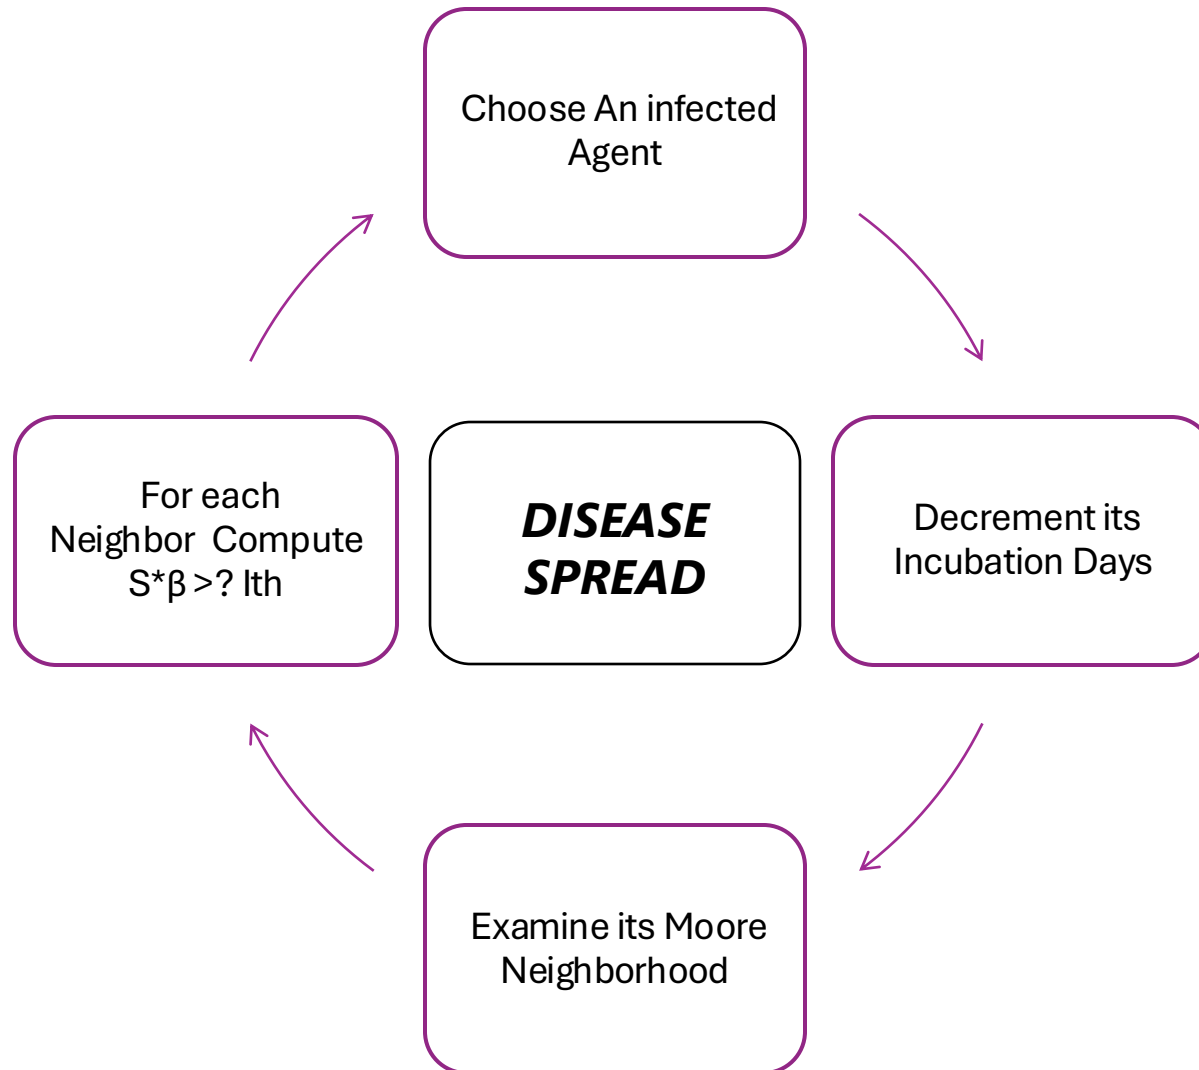
Solution Description: OpenMP Version – Daily Simulation




Solution Description: OpenMP Version – Daily Simulation




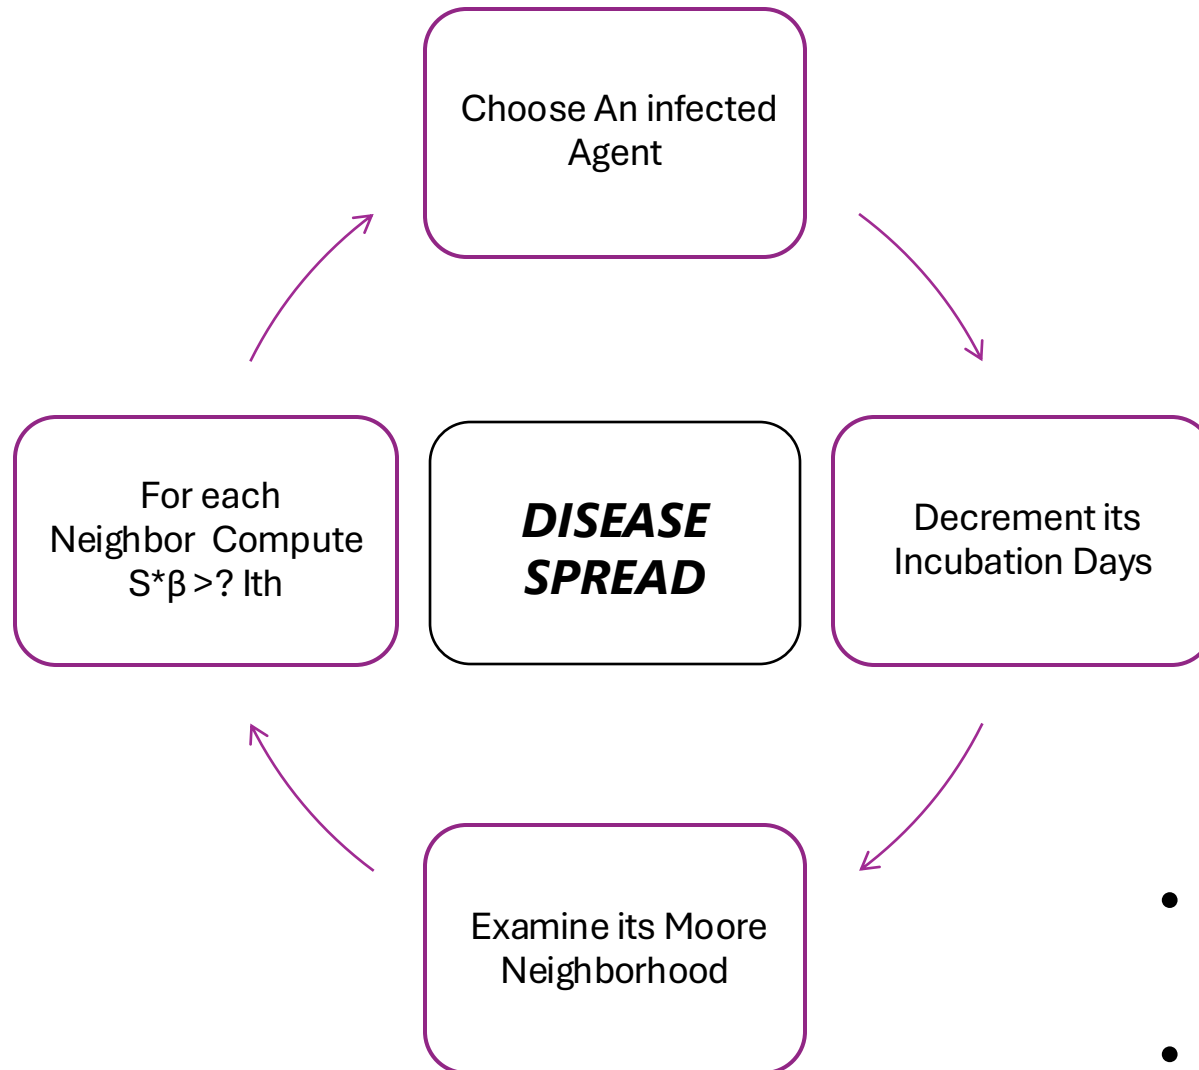
Solution Description: OpenMP Version – Daily Simulation



TID 1





Solution Description: OpenMP Version – Daily Simulation




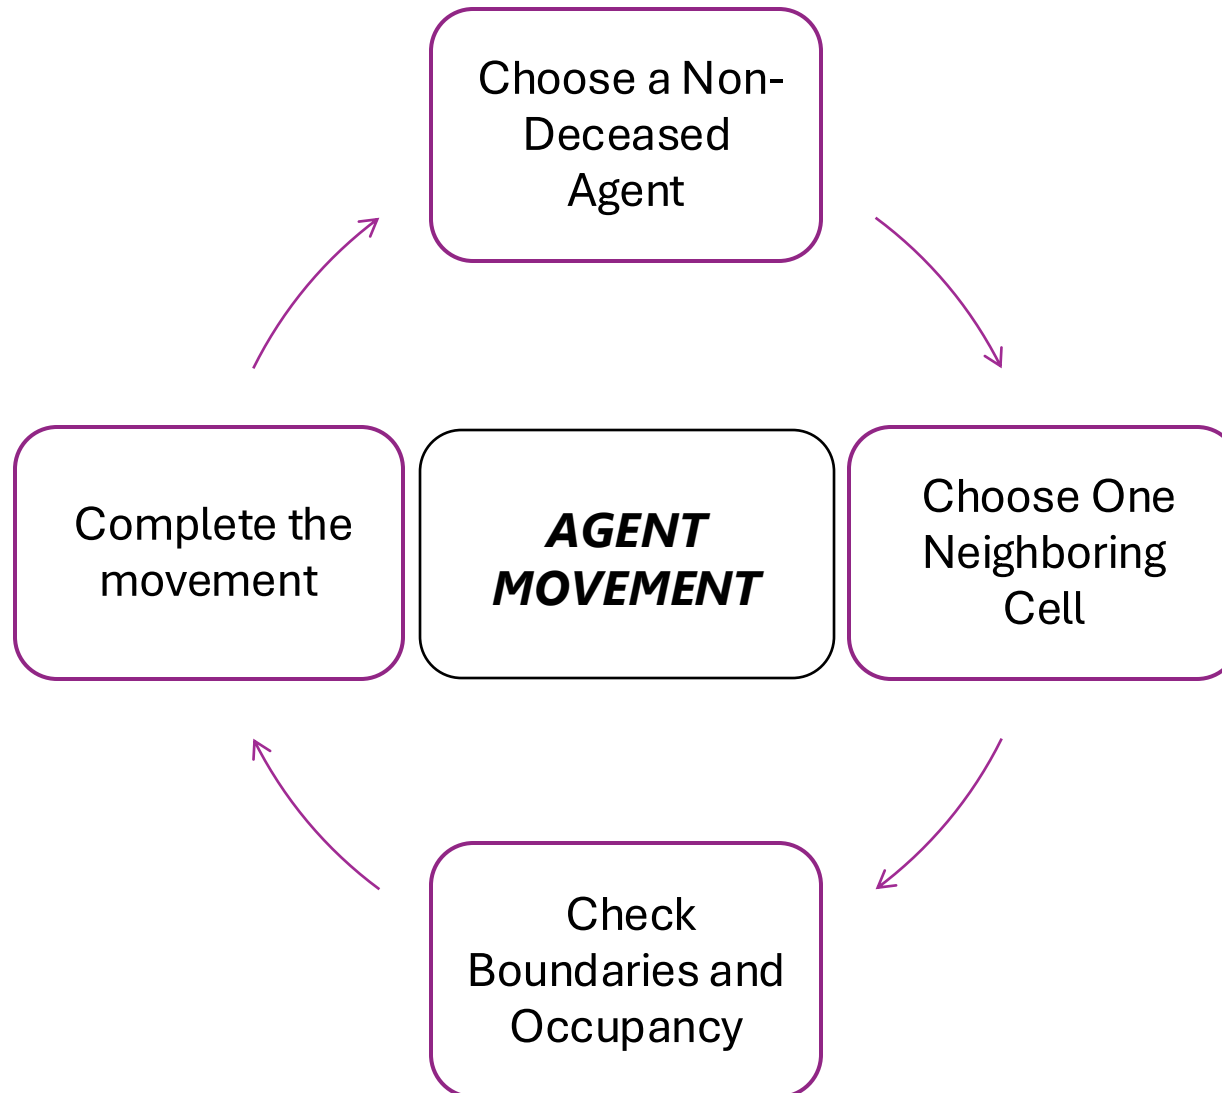
- **Per-Thread Newly infected**
- ***incubation_days* = -1**




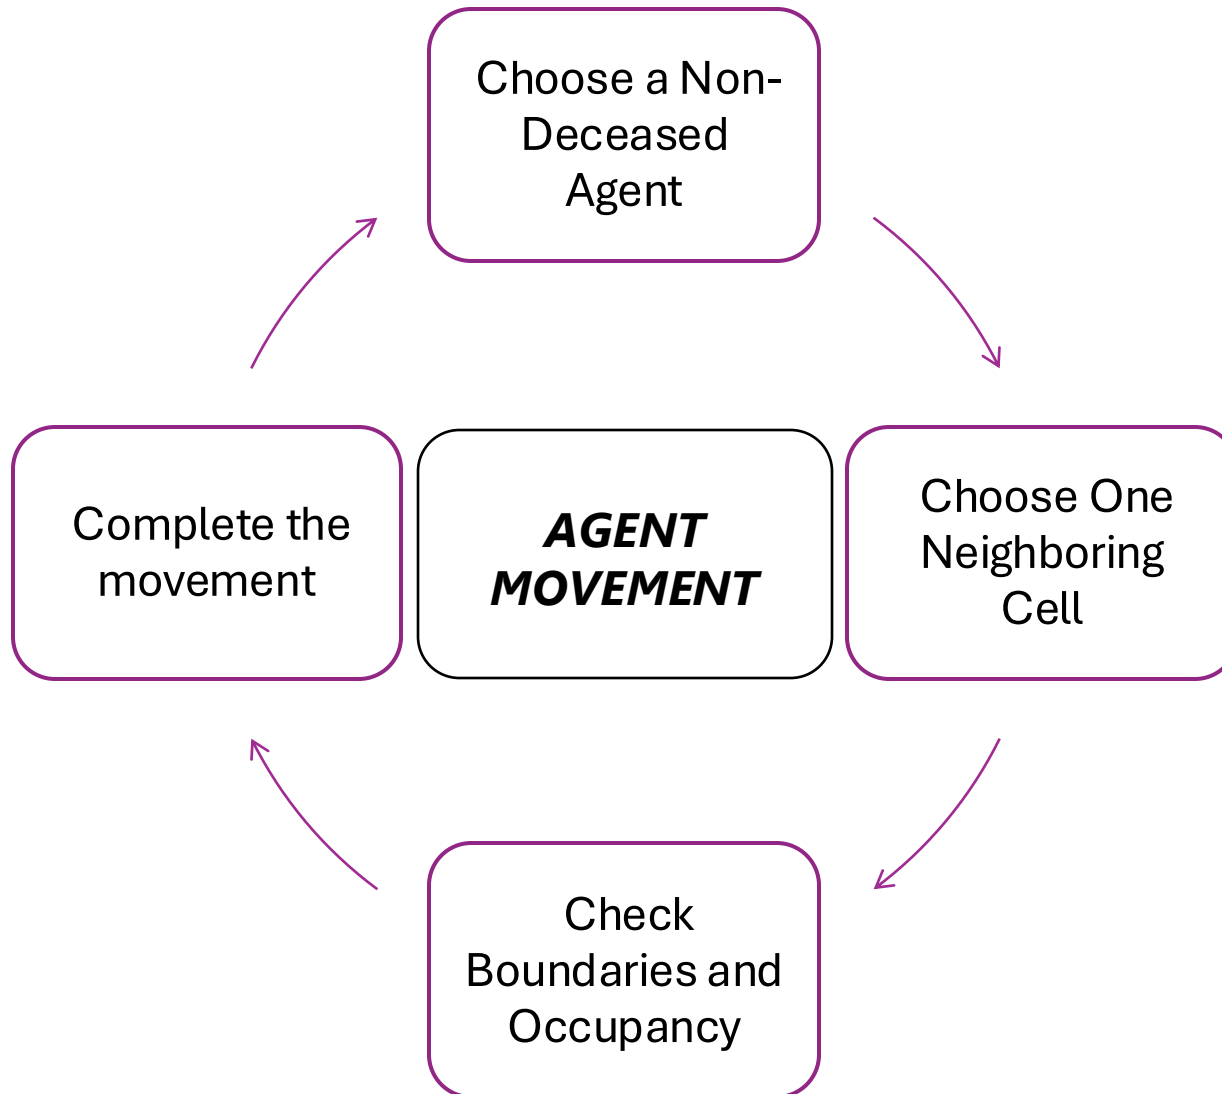
Solution Description: OpenMP Version – Daily Simulation




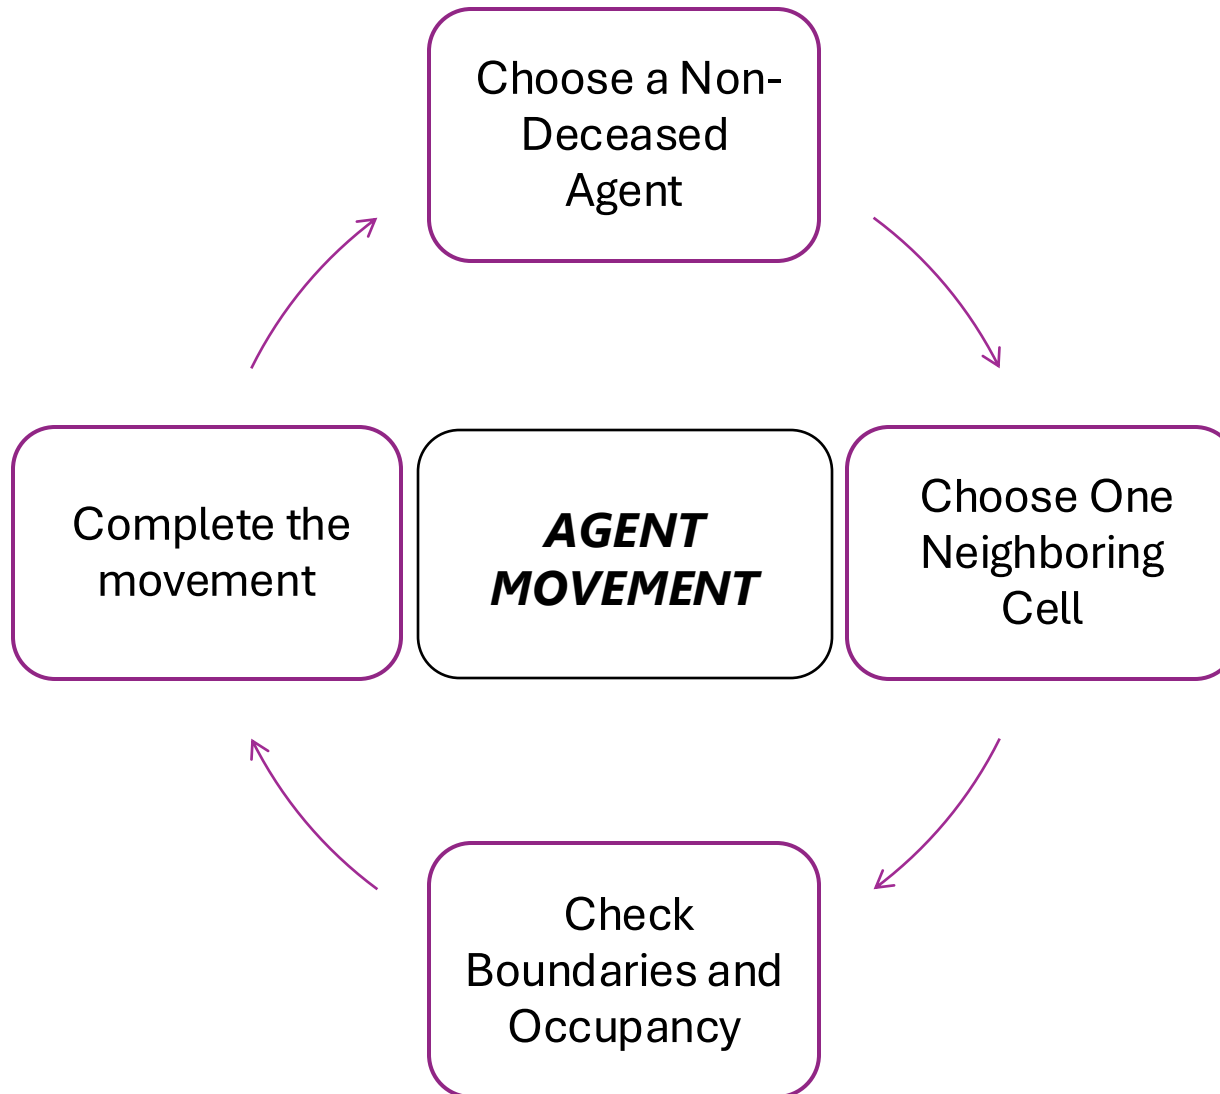
Solution Description: OpenMP Version – Daily Simulation




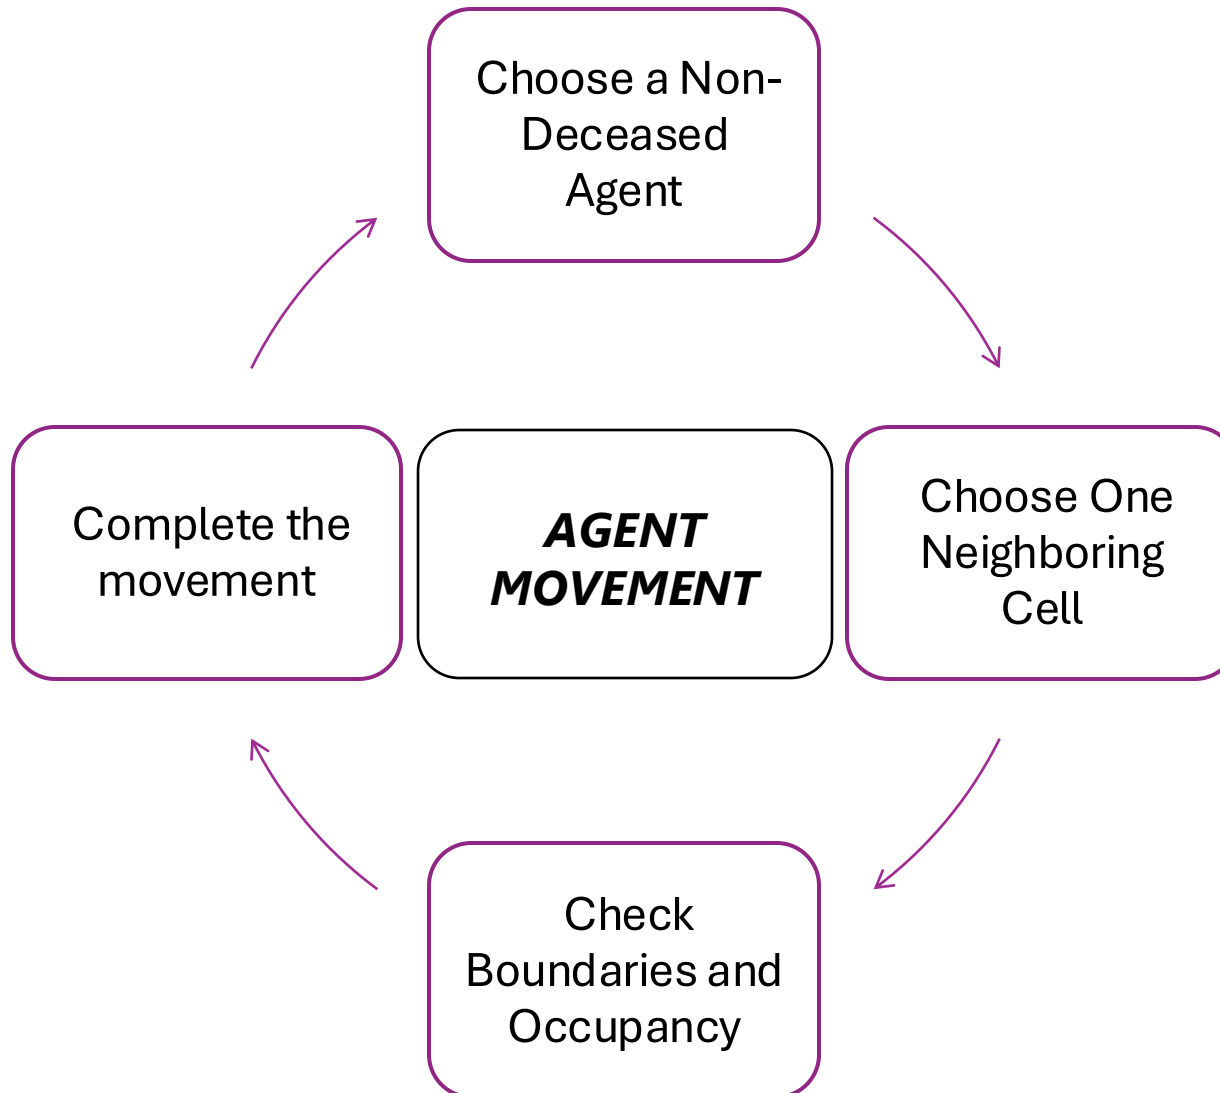
Solution Description: OpenMP Version – Daily Simulation





Solution Description: OpenMP Version – Daily Simulation




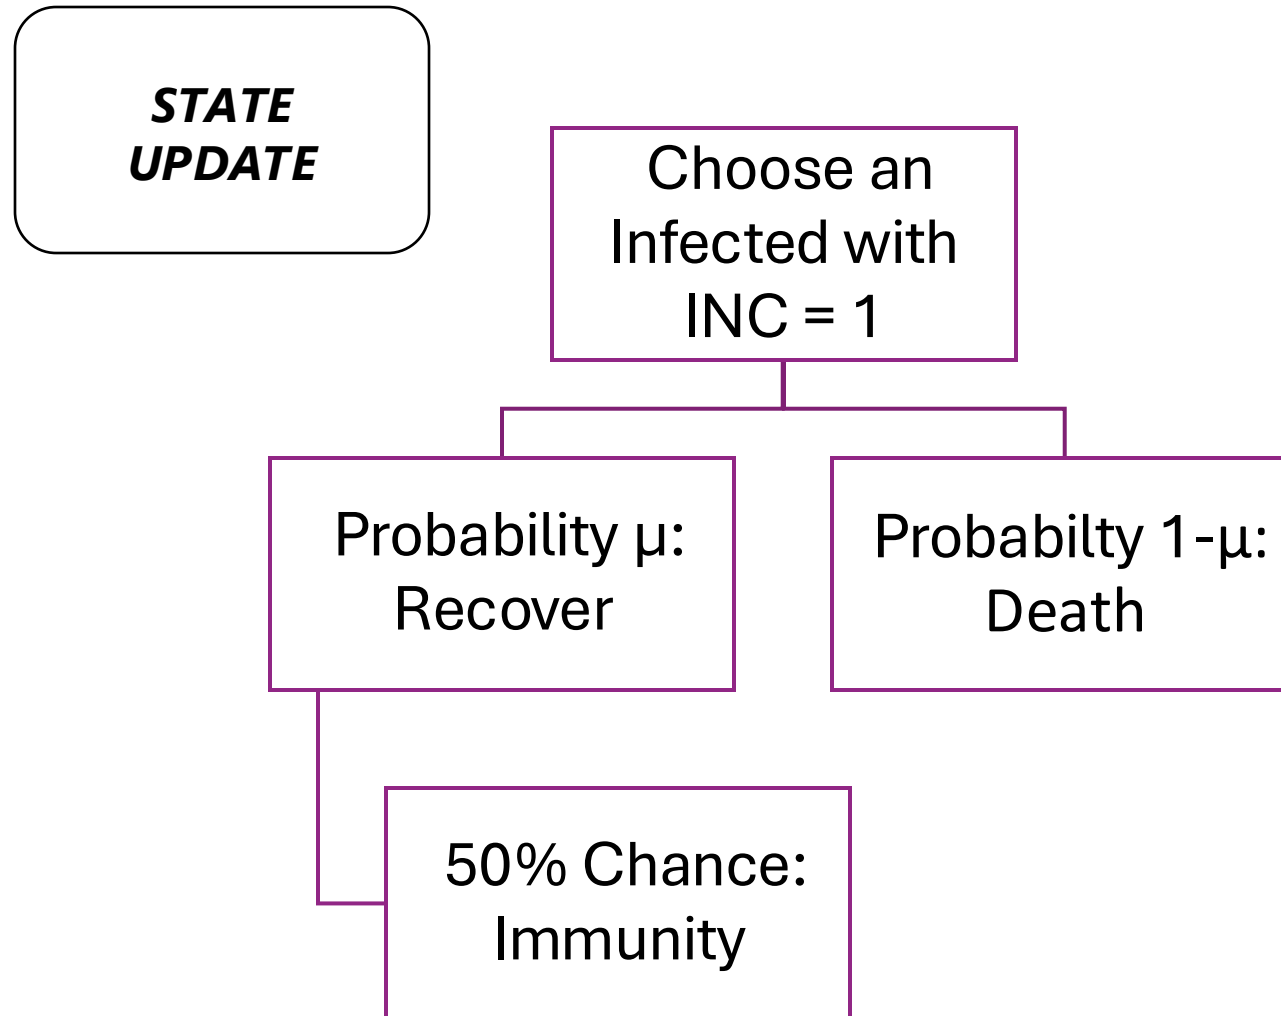
Solution Description: OpenMP Version – Daily Simulation



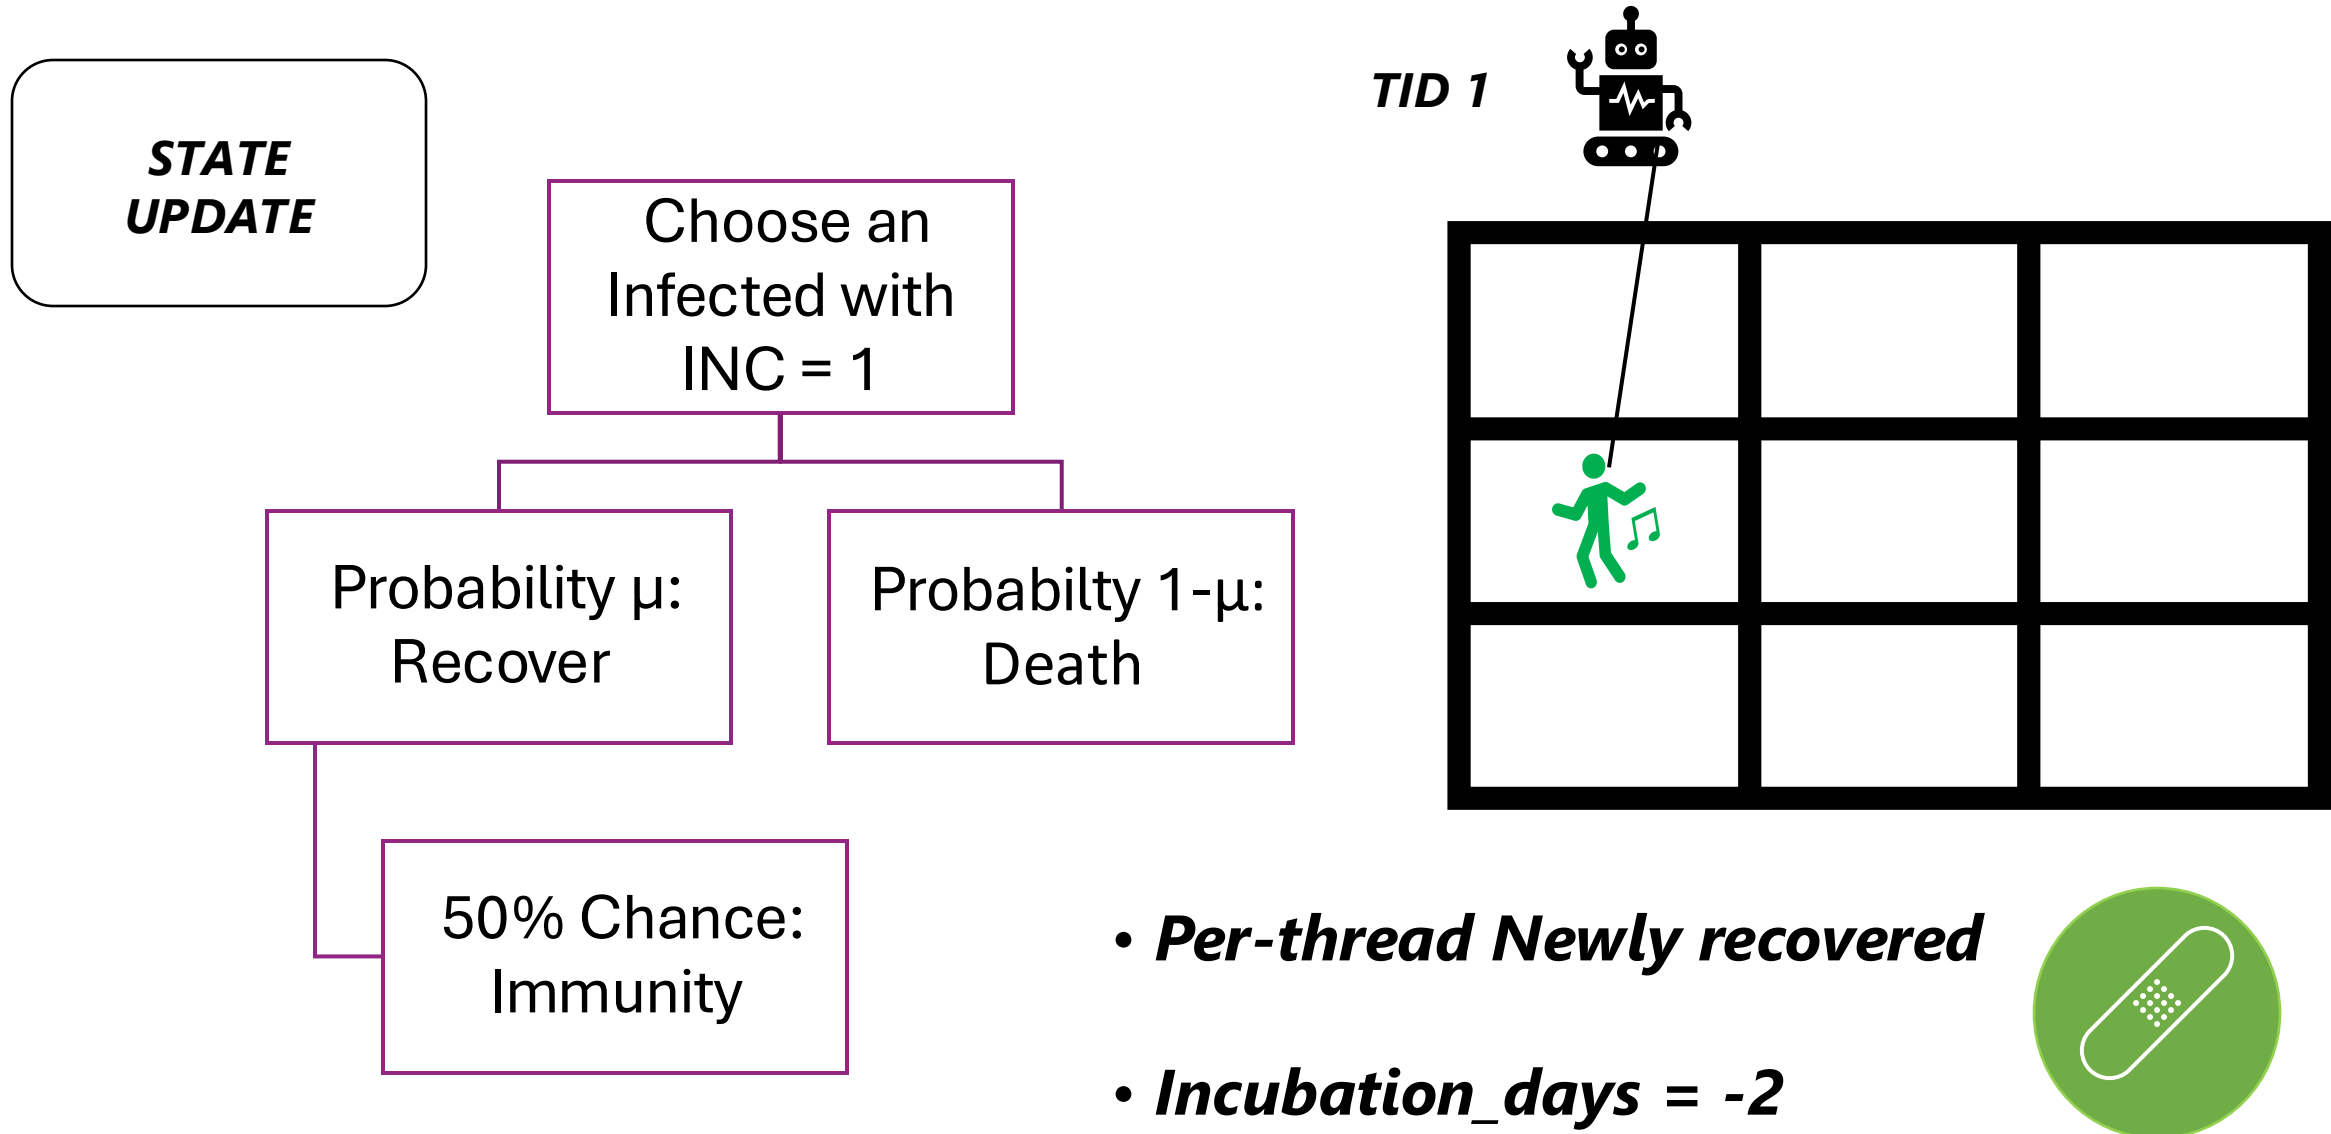
Solution Description: OpenMP Version – Daily Simulation




Solution Description: SERIAL Version – Daily Simulation



Solution Description: SERIAL Version – Daily Simulation

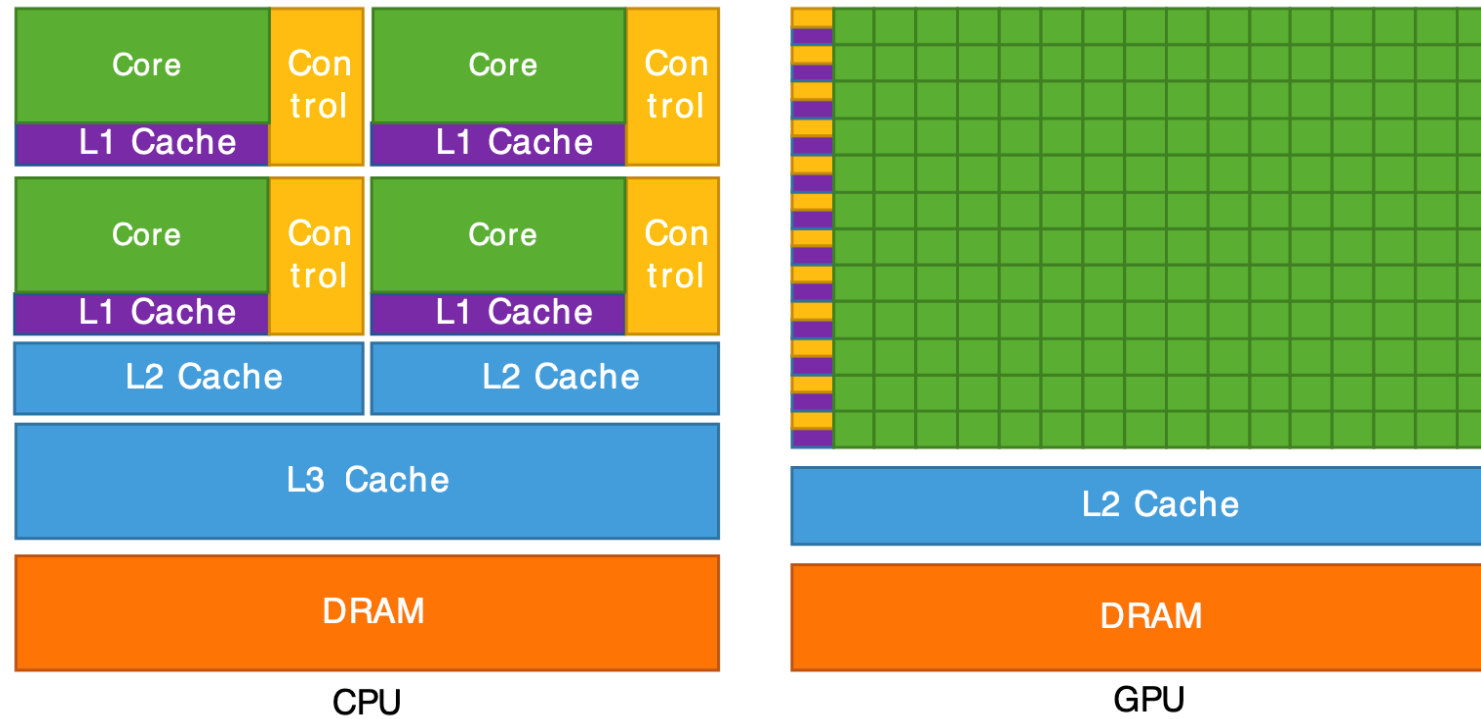


Solution Description: OpenMP Version – Daily Simulation



Solution Description: CUDA Version

Change of architecture



Solution Description: CUDA Version

Thread



=






Person

```
int threads = # multiple of a warp  
int blocks = (NP + threads - 1) / threads;
```





Grid = Population

Solution Description: CUDA Version



Tiling



Global synchronization

Solution Description: CUDA Version --> Togliere codice, rimettere le icone di slide 6

coalesced memory access

Array Of Structures

```
typedef struct {  
    int x, y;  
    float susceptibility;  
    int incubation_days;  
} Person;  
  
typedef struct {  
    int occupancy;  
    Person **persons;  
} Cell;
```



Structure of Arrays

```
int *d_x;  
int *d_y;  
int *d_incub;  
int *d_newInf;  
float *d_susc;  
  
int *d_cellCount;  
int *d_slotIndex;  
int *d_cellSlots;
```

Solution Description: CUDA Version

cell count (W*H)

1	0	2	1
---	---	---	---


slot index (NP)

0	0	1	0
---	---	---	---

cell slots (W*H*MaxP)

0	#		#	#	1	2	3	#
---	---	--	---	---	---	---	---	---

Max People per cell = 2



Solution Description: CUDA Version

cell count ($W \times H$)


1	0	2	1
---	---	---	---

slot index (NP)

0	0	1	0
---	---	---	---

cell slots ($W \times H \times \text{MaxP}$)

0	#	#	#	1	2	3	#
---	---	---	---	---	---	---	---



Solution Description: CUDA Version

cell count ($W \times H$)


1	0	2	1
---	---	---	---

slot index (NP)

0	0	1	0
---	---	---	---

cell slots ($W \times H \times \text{MaxP}$)

0	#	#	#	1	2	3	#
---	---	---	---	---	---	---	---



Solution Description: CUDA Version

cell count ($W \times H$)


1	0	2	1
---	---	---	---

slot index (NP)

0	0	1	0
---	---	---	---

cell slots ($W \times H \times \text{MaxP}$)

0	#	#	#	1	2	3	#
---	---	---	---	---	---	---	---



Solution Description: CUDA Version

cell count (W*H)

1	0	2	1
---	---	---	---





slot index (NP)

0	0	1	0
---	---	---	---

old Slot index

0

Max People per cell = 2

0 	1 
2 	3 

cell slots (W*H*MaxP)

0	#	1	#	1	2	3	#
---	---	---	---	---	---	---	---

Solution Description: CUDA Version

cell count (W*H)

1	1	1	1
---	---	---	---


slot index (NP)

0	0	0	0
---	---	---	---

old Slot index

0

Max People per cell = 2



cell slots (W*H*MaxP)

0	#	1	#	1	2	3	#
---	---	---	---	---	---	---	---

Solution Description: CUDA Version

*cell count (W*H)*

1	1	1	1
---	---	---	---





slot index (NP)

0	0	0	0
---	---	---	---

old Slot index

0

Max People per cell = 2

0 	1 
2 	3 


*cell slots (W*H*MaxP)*

0	#	1	#	2	#	3	#
---	---	---	---	---	---	---	---


Solution Description: CUDA Version



infection kernel



*status update
kernel*



movement kernel


Solution Description: CUDA Version



Infection kernel



```
if (infec > ITH)
{
    atomicCAS(&d_newInf[i], 0, 1);
}
```



*status update
kernel*



movement kernel


Solution Description: CUDA Version



Infection kernel



```
if (infec > ITH)
{
    atomicCAS(&d_newInf[i], 0, 1);
}
```



*status update
kernel*



movement kernel



```
int pos = atomicAdd(&d_cellCount[newC], 1);
if (pos < MAXP_CELL)
{
    ...
}
else
{
    // rollback
    atomicSub(&d_cellCount[newC], 1);
}
```

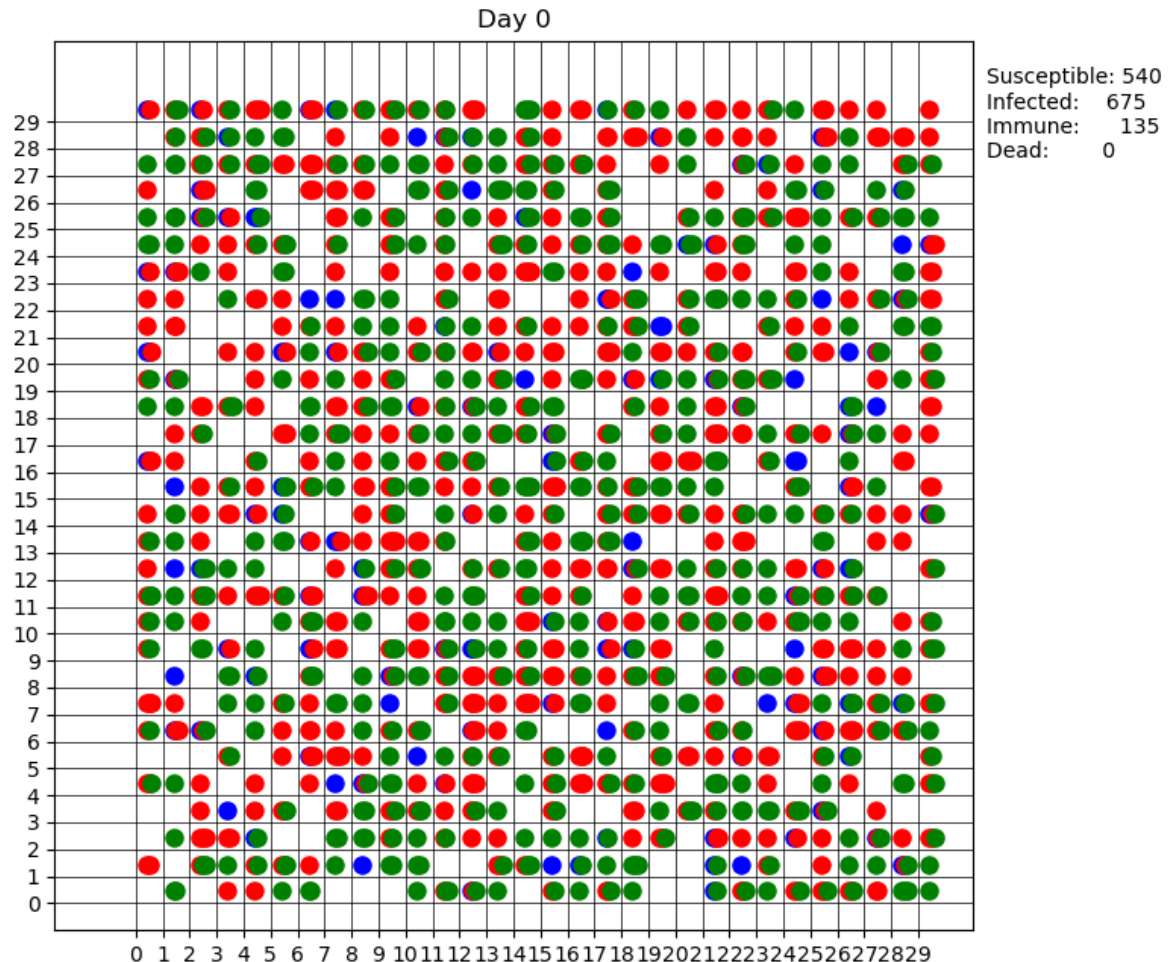
Solution Description: CUDA Version

***all data structures are on
device except for***

```
clock_gettime(CLOCK_MONOTONIC, &ts_rm1);  
cudaMemcpy(d_x, h_x, NP * sizeof(int), cudaMemcpyHostToDevice);  
cudaMemcpy(d_y, h_y, NP * sizeof(int), cudaMemcpyHostToDevice);  
clock_gettime(CLOCK_MONOTONIC, &te_rm1);
```

***we expect
 $Cuda \cong Cuda\ Evo$***

Solution Description: Report



- *Report of the simulation for each day*
- *Saving data about state and positions*
- *Saving data to binary files:
implementation-agnostic manner*

Solution Description: Report

Unit test

cell occupancy limit

deaths never decrease

initial immune and infected counts

incubation duration and transition

cell movement

state counts sum to NP

positions valid or dead

entries per day

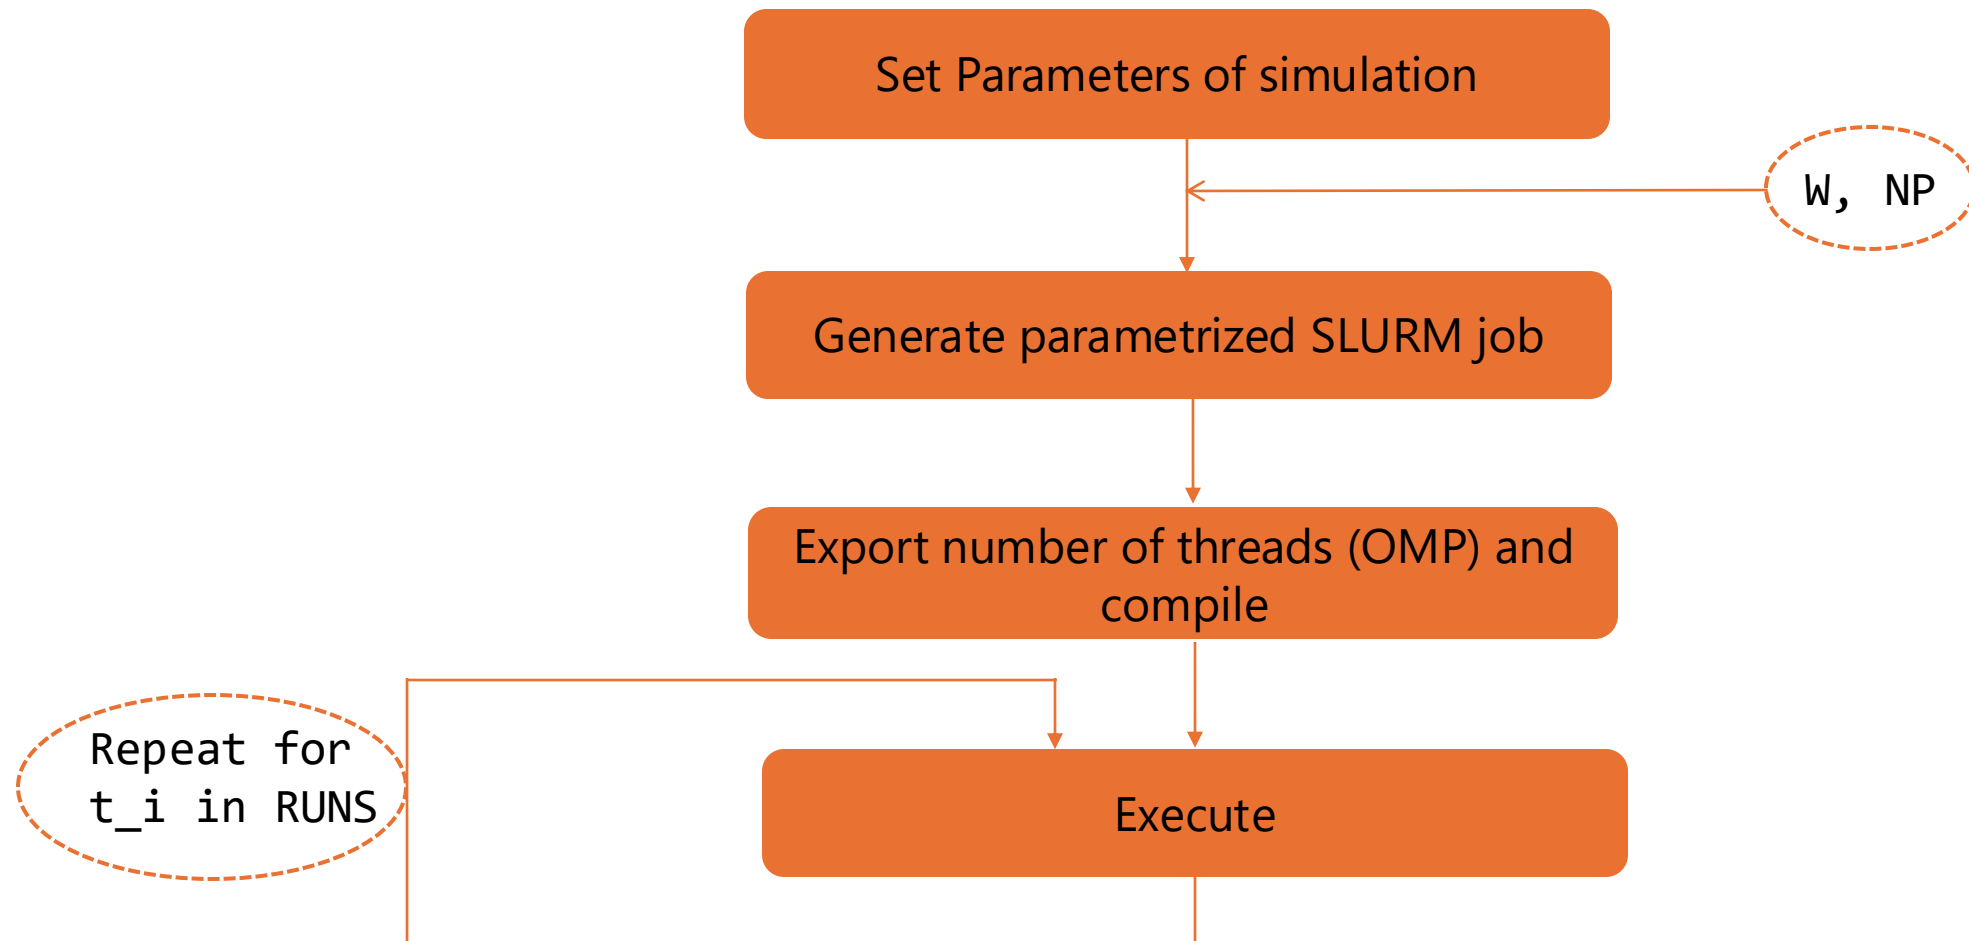
Experimental Results: Platform

- ***HPC Polito (Hactar)***
 - ***CPU Node:*** 2x Xeon E5-2680 v3 2.50 GHz (turbo 3.3 GHz) 12 cores
 - ***GPU Node:*** 2x Tesla K40 – 12GB – 2880 cuda cores
 - ***Job Scheduler:*** SLURM




Experimental Results: Simulation File Generation

- Study **performance** of the three implementations with increasing input dimension: Serial, OpenMP and Cuda



Experimental Results: Performance Analysis



DATASET:

- RUNS: 10
- Logical parameters: default values
- Grid Sizes: from 10x10 to 10000x10000
- Population Densities: 0.1, 0.5 and **0.9**


OPENMP:

- Number of Threads = {2, 4, 8, 16, 24}

CUDA:

- Number of Threads per Block = {128, 256, 1024}


Experimental Results: Speedup - OpenMP



SPEEDUP $\in [2, 24]$

- $\frac{t_{serial}}{t_{OMP(N)}} \approx N$, with N number of threads

Experimental Results: Speedup - OpenMP



SPEEDUP $\in [2, 24]$


- $\frac{t_{serial}}{t_{OMP(N)}} \approx N$, with N number of threads

- **Linear trend** in line with Amdahl Law:

$$t(N) = t_{serial} + \frac{t_{parallel}}{N}$$

- $t_{serial} \ll t_{parallel}$
- Slight fluctuations due to inherent randomness


Experimental Results: Speedup - CUDA



SPEEDUP $\in [55, 80]$

- Huge performance **gain** for large grids thanks to thousands of threads
- Speedup almost **independent** of number of threads per block

Experimental Results: Speedup - CUDA




SPEEDUP $\in [55, 80]$


- Huge performance **gain** for large grids thanks to thousands of threads
- Speedup almost **independent** of number of threads per block
- **Negligible impact** of data transfers on performance (CUDA \approx CUDA EVO)

Experimental Results: Logical Simulation


- **Q:** Most **critical** parameter for population?
 - **A:** Fix initial conditions and simulate evolution of population for varying S_{AVG} and μ



$$S_{AVG} = 0.3, \mu = 0.6$$



$$S_{AVG} = 0.6, \mu = 0.6$$




$$S_{AVG} = 0.9, \mu = 0.6$$

→ Threshold – behaviour vs S_{AVG} due to infection condition:


$$S_i >? \frac{ITH}{\beta}$$

Experimental Results: Logical Simulation


- **Q:** Most **critical** parameter for population?
 - **A:** Fix initial conditions and simulate evolution of population for varying S_{AVG} and μ



$$S_{AVG} = 0.7, \mu = 0.9$$



$$S_{AVG} = 0.7, \mu = 0.6$$



$$S_{AVG} = 0.7, \mu = 0.3$$

→ By the Law of Large numbers:

$$num_Immune \approx \mu \times num_Dead \quad \text{as } t \rightarrow \infty$$

Conclusions

- **Problem:** Simulate propagation of epistemics in a population
- **Three solutions:** Serial, OpenMP and Cuda:
 - Testing
 - Performance Evaluation on HPCPolito
 - Logical Simulation Analysis
- **Parallel Speedup:**
 - **OMP:** $\in [2, 24]$ linear with the number of threads
 - **CUDA:** ≈ 60 , approx. constant with number of threads per block, no impact of memory transfer

Thank you for your attention!

Federico Muscarà Marco Parentin Leonardo Straccali Francesco Zanasi